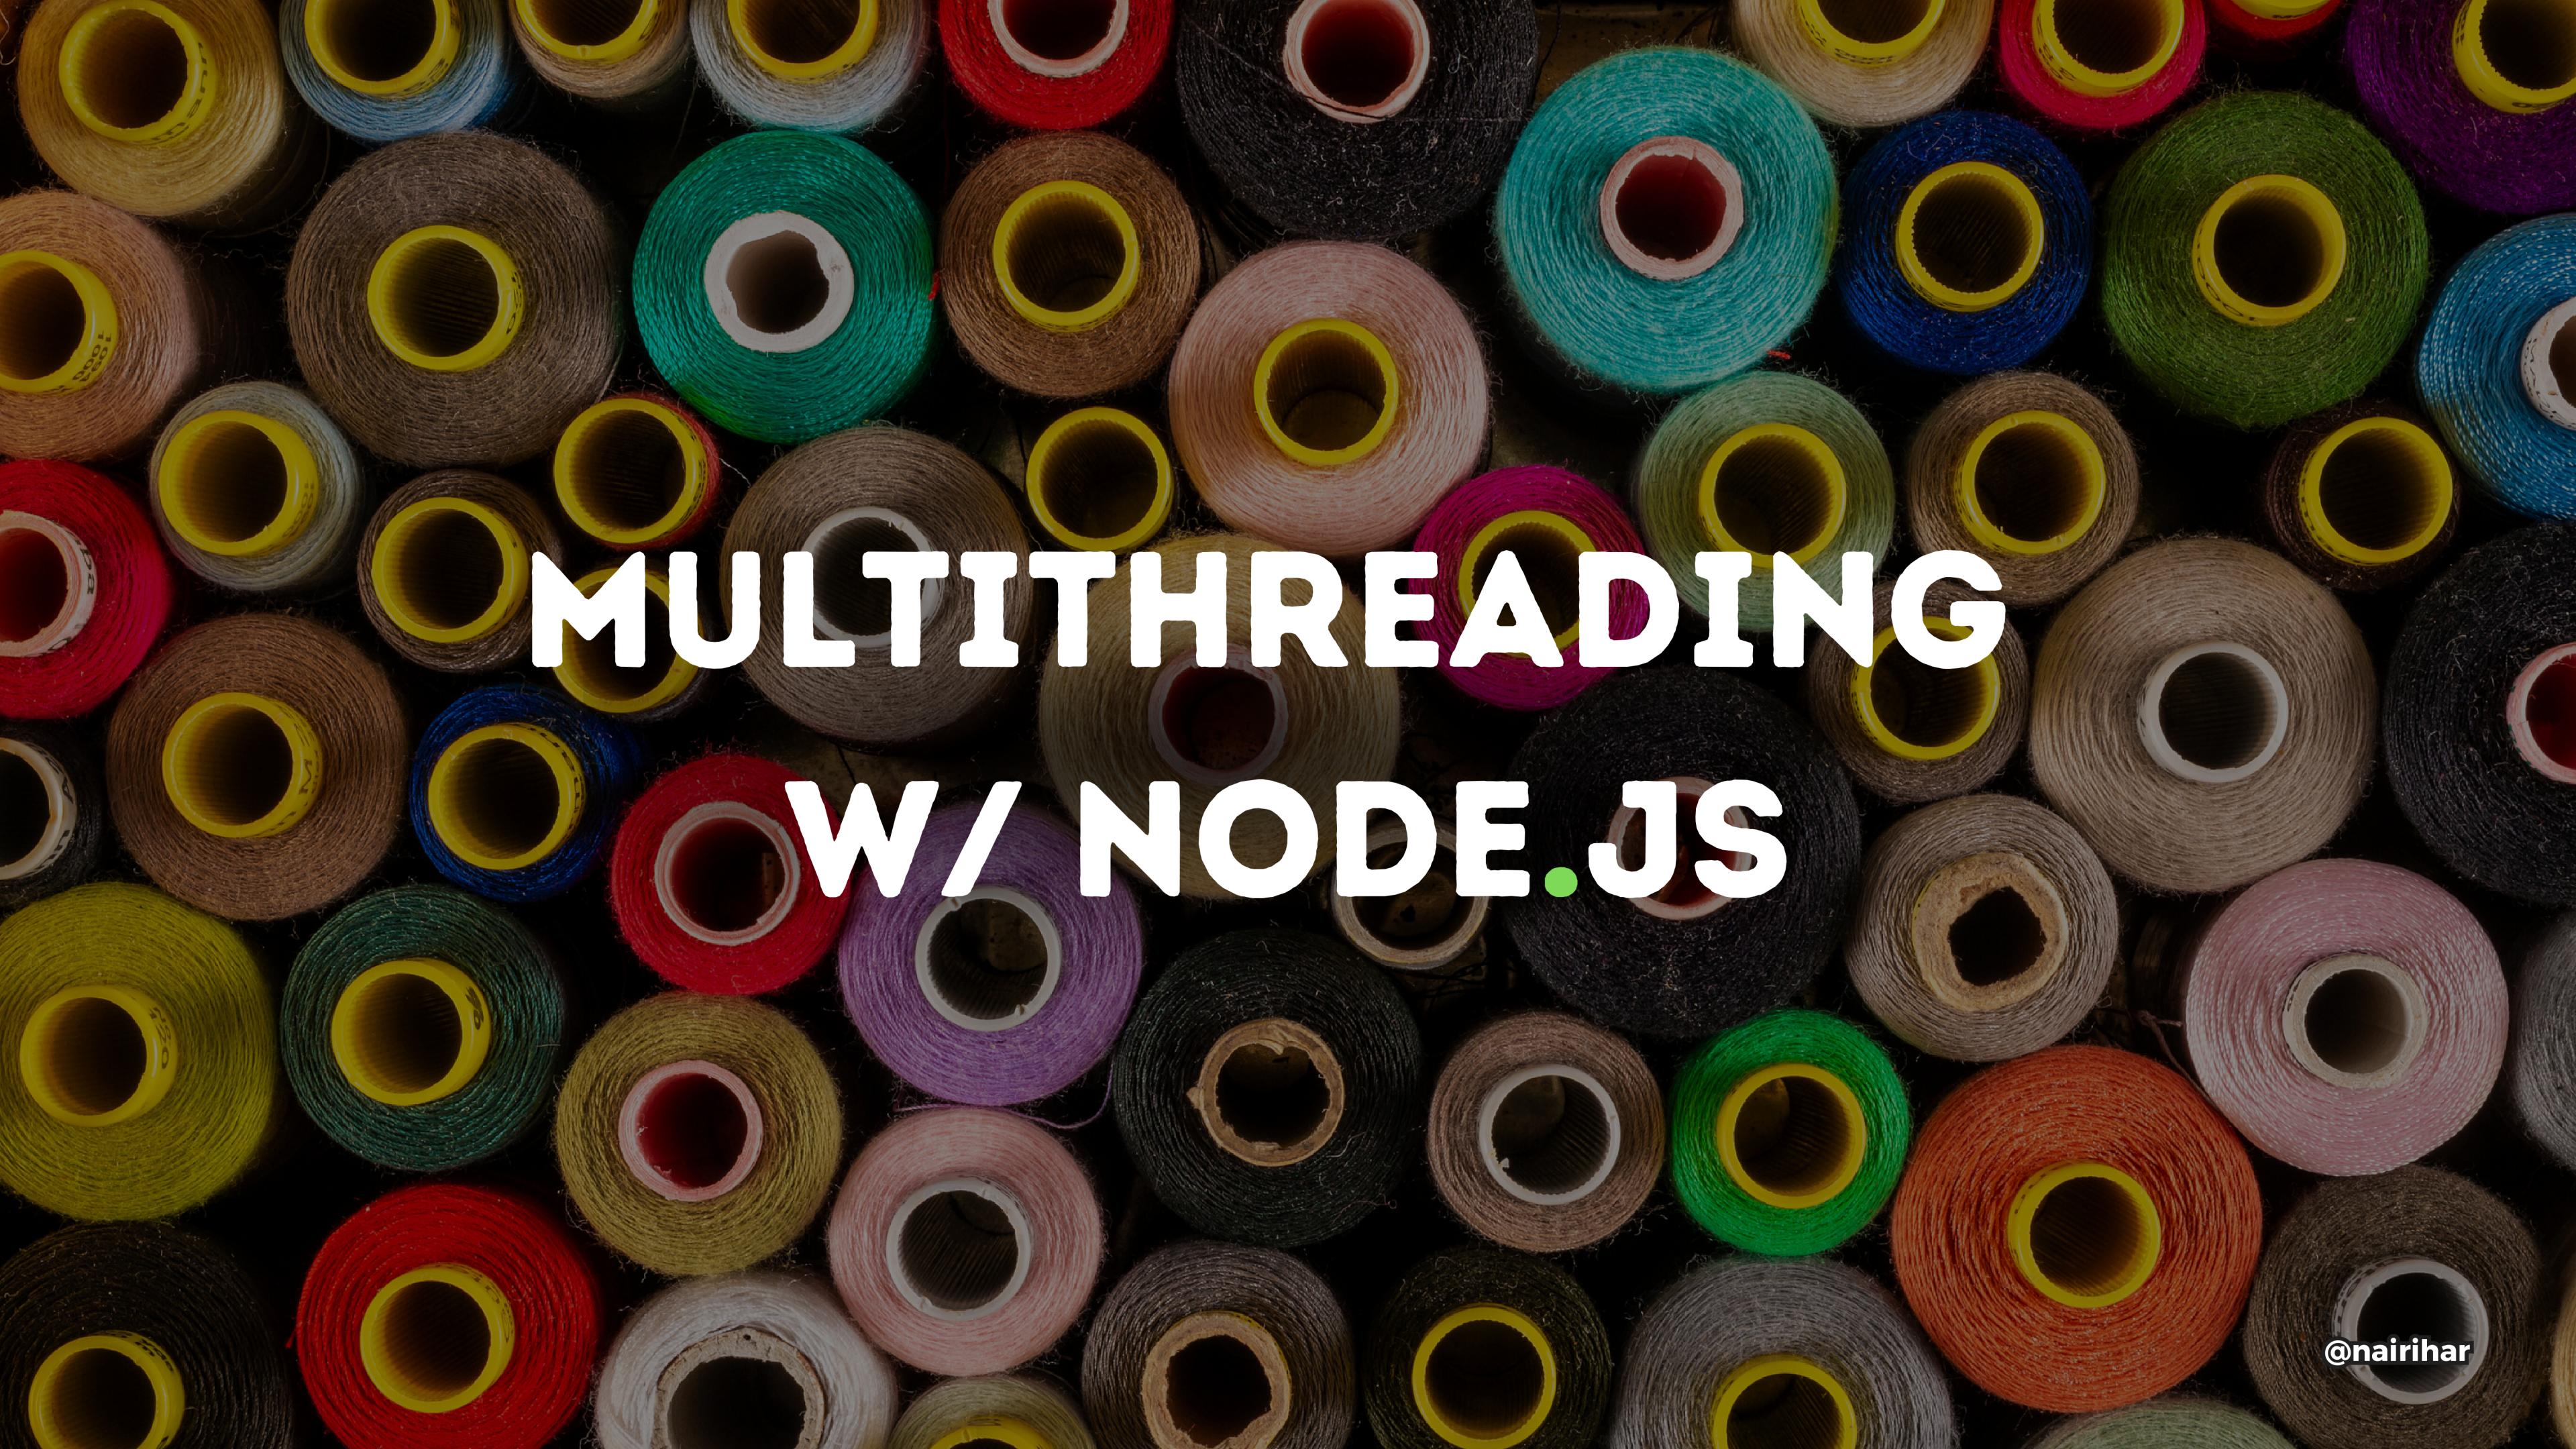




@nairihar



MULTITHREADING w/ NODE.JS

Before a brief introduction
about myself . . .

HEY I'M NAIRI *_*

- 💻 Backend Engineer at Screenful
- ❤️ Creator of the **JavaScript Armenia** Community
- ✍️ Writing about JavaScript and related technologies
- 🎙️ Hosting podcasts for the Armenian tech community
- 🎤 Speaker, Conference/Meetup Organizer



@nairihar

Do you already know
about **worker_threads**?

AGENDA

- Node.js & Libuv (not so single-threaded)
- worker_threads (basics + demo)
- SharedArrayBuffer (shared data)
- Atomics (race conditions)
- Abstractions for Node.js Multithreading

JavaScript



Simpler than you think...

We use it everywhere...

WEB

MOBILE

IOT

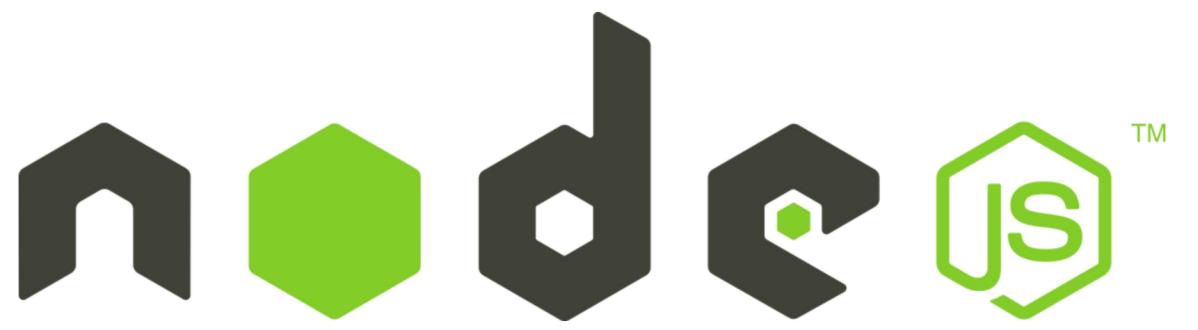
BACKEND

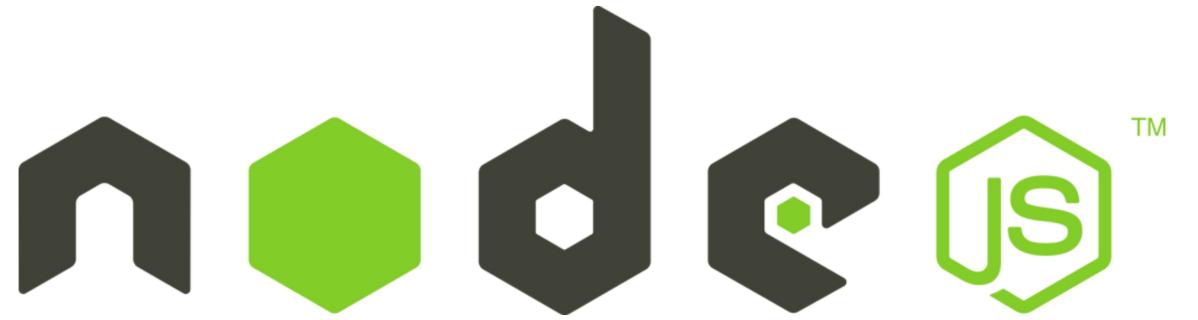


DESKTOP

GAME

...





libuv

I/O in Operation systems

I/O in Operation systems

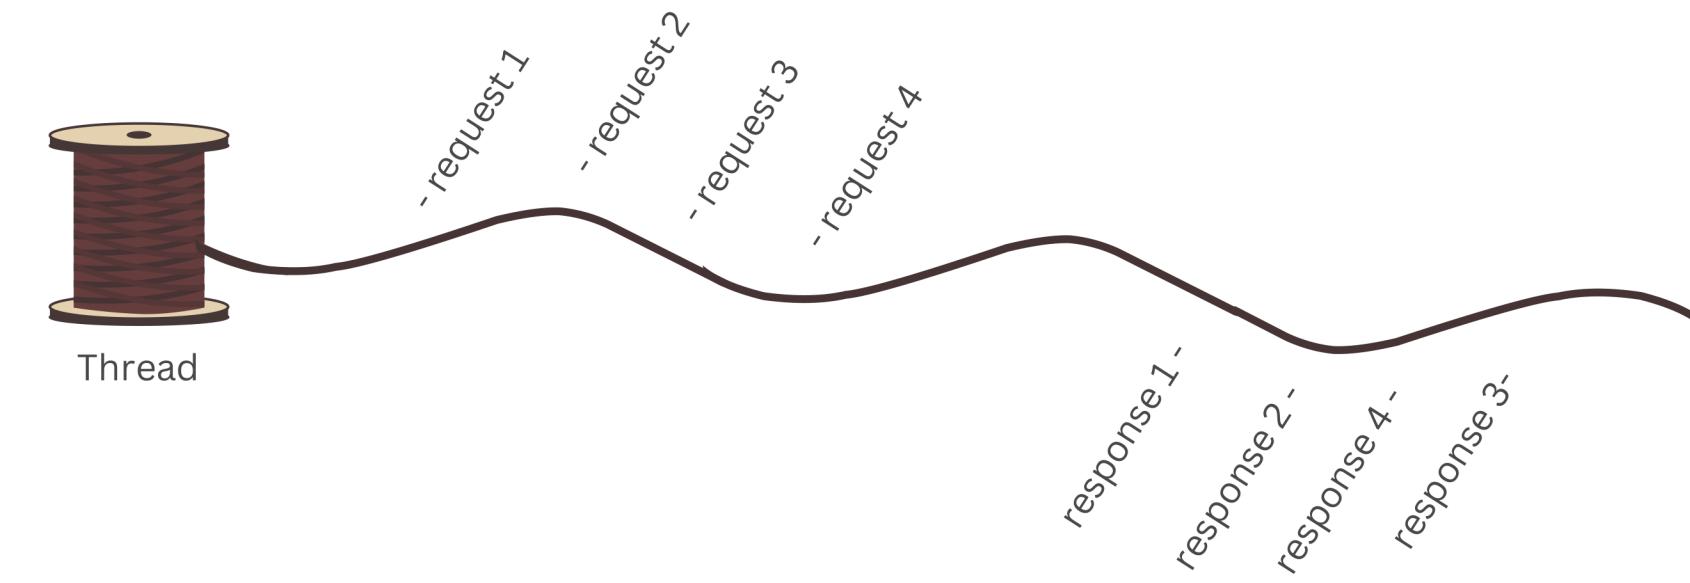
Non blocking I/O

- *Network*

I/O in Operation systems

Non blocking I/O

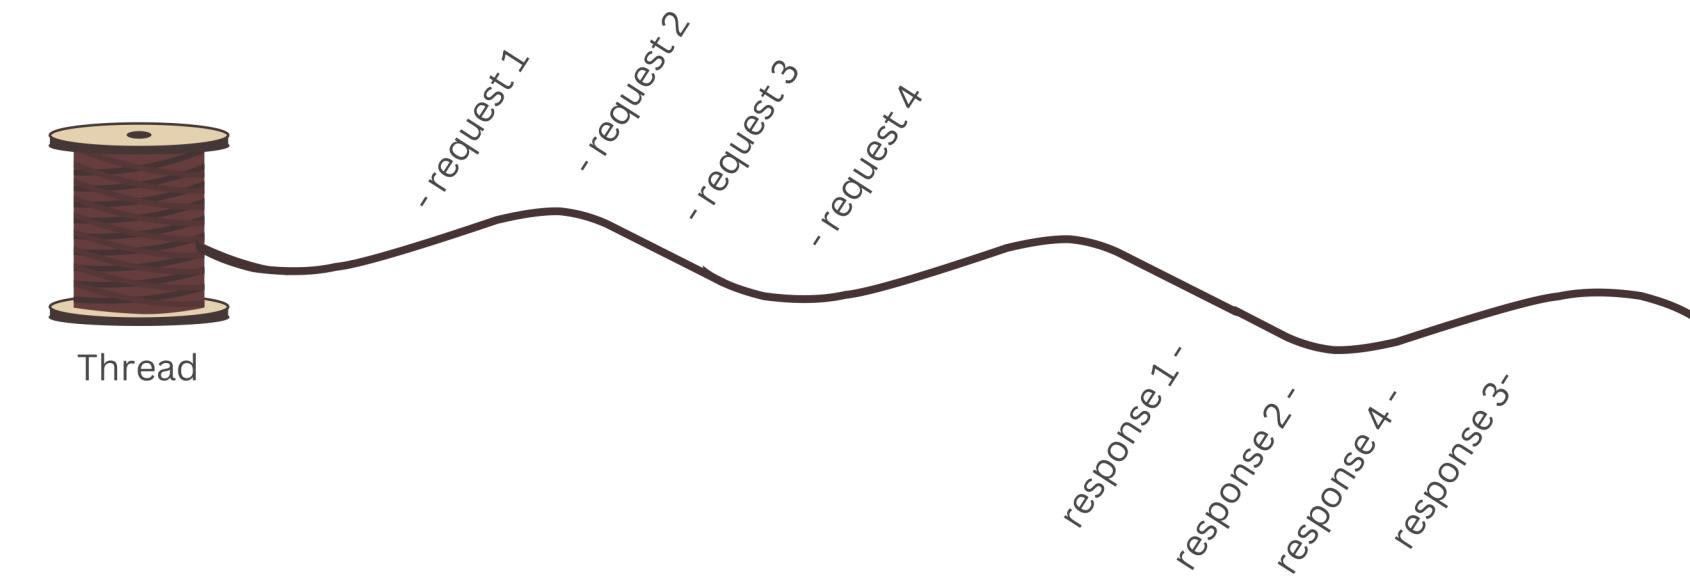
- *Network*



I/O in Operation systems

Non blocking I/O

- *Network*



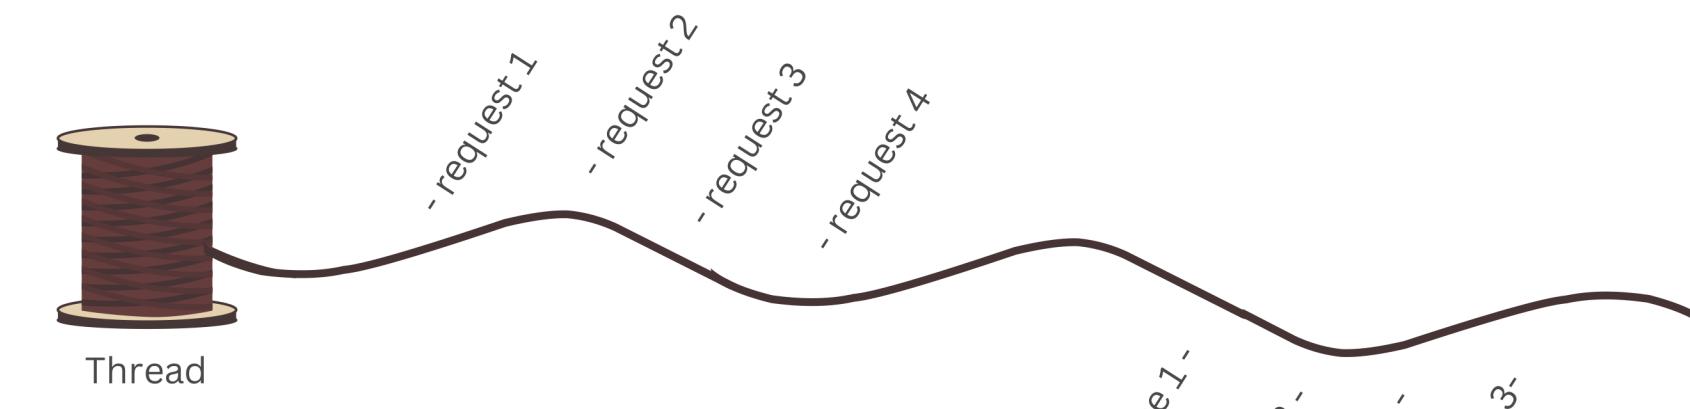
Blocking I/O

- *File system*
- *DNS.lookup*

I/O in Operation systems

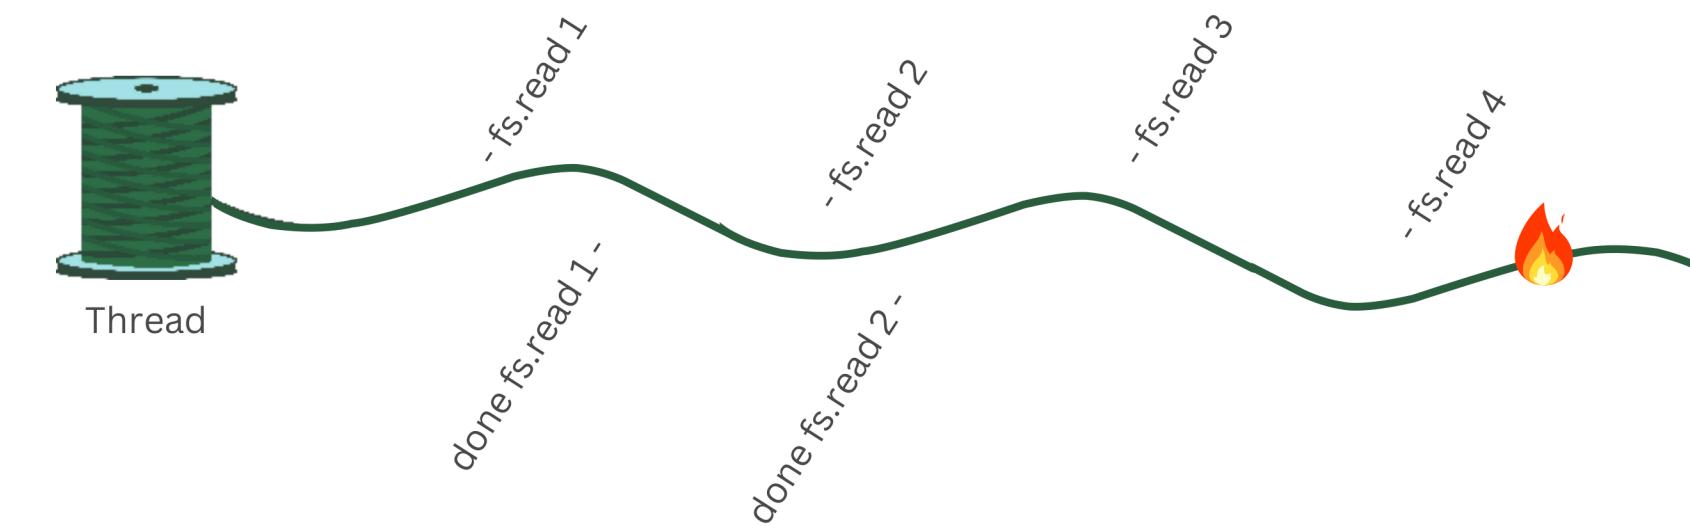
Non blocking I/O

- *Network*



Blocking I/O

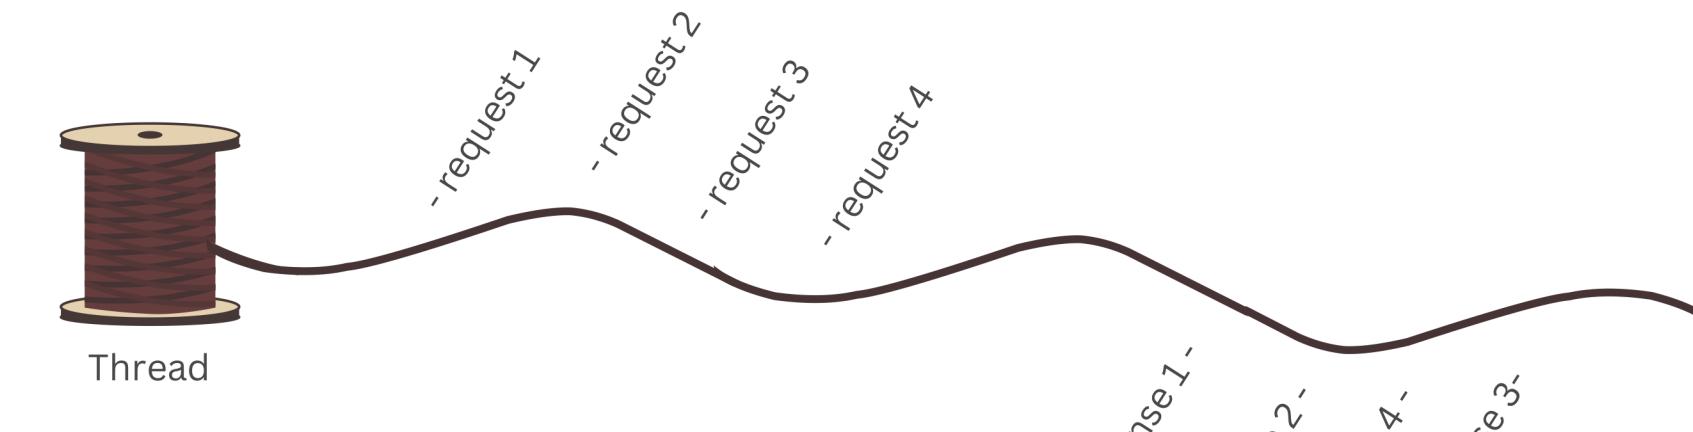
- *File system*
- *DNS.lookup*



I/O in Operation systems

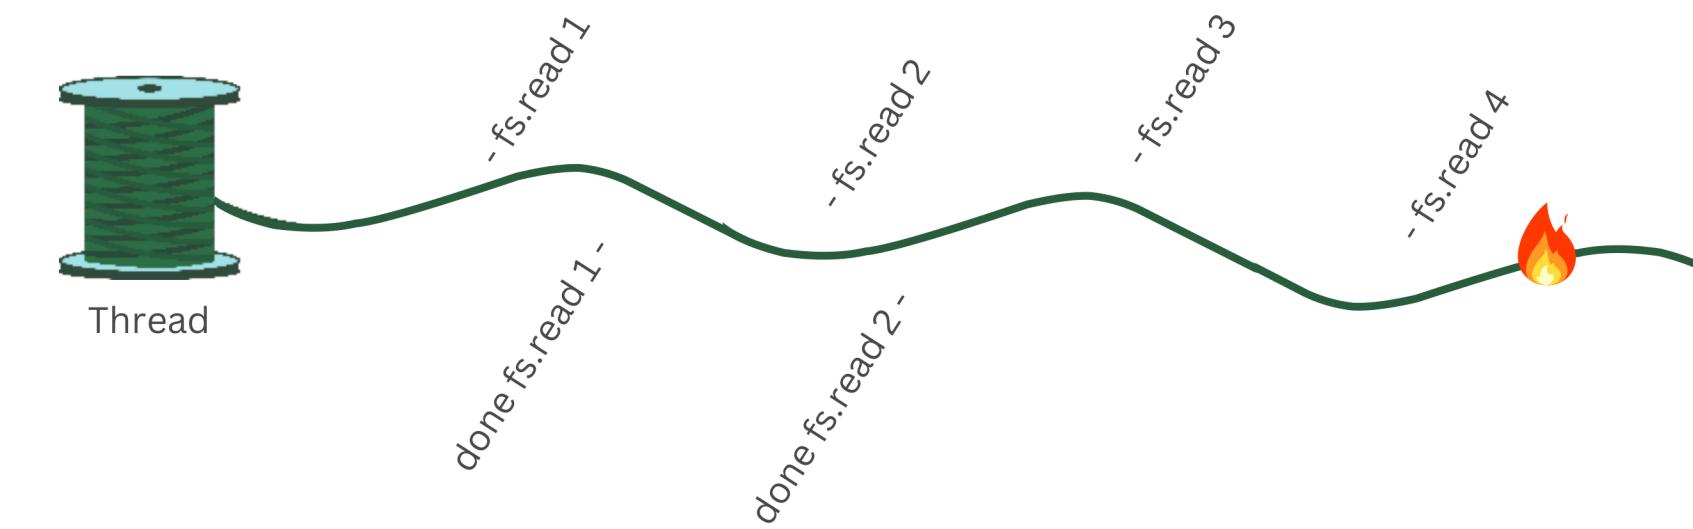
Non blocking I/O

- *Network*



Blocking I/O

- *File system*
- *DNS.lookup*



***io_uring** is a Linux kernel system call interface for storage device **asynchronous I/O** operations.

@nairihar

Node.js is a single-threaded runtime

```
const fs = require('fs').promises;

(async () => {
  const [ user, balance ] = await Promise.all([
    fs.readFile('user.json'),
    fs.readFile('balance.json'),
  ]);

  // ...
})()
```



How?

Blocking I/O & CPU intensive

File I/O

DNS

...

Thread Pool



Thread Pool



Thread Pool



UV_THREADPOOL_SIZE=10 (Up to 1024 threads)

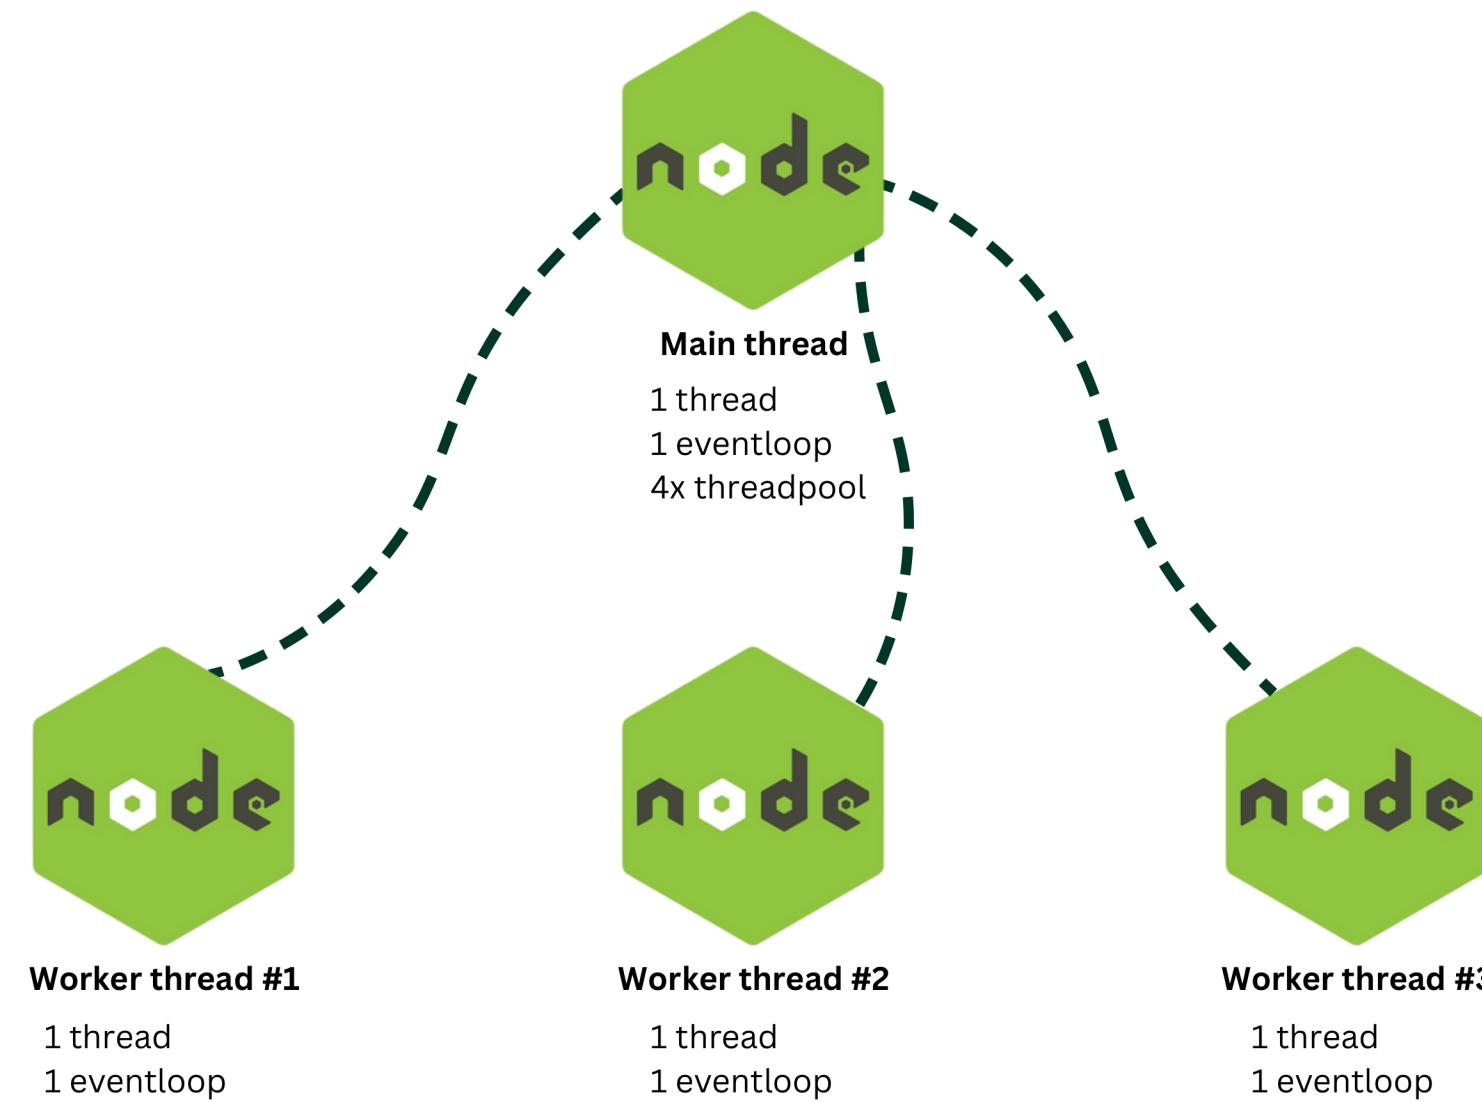
Before node v12.5.0: Up to 128 threads (before libuv v1.30.0)

We can't use **ThreadPool**
threads directly...

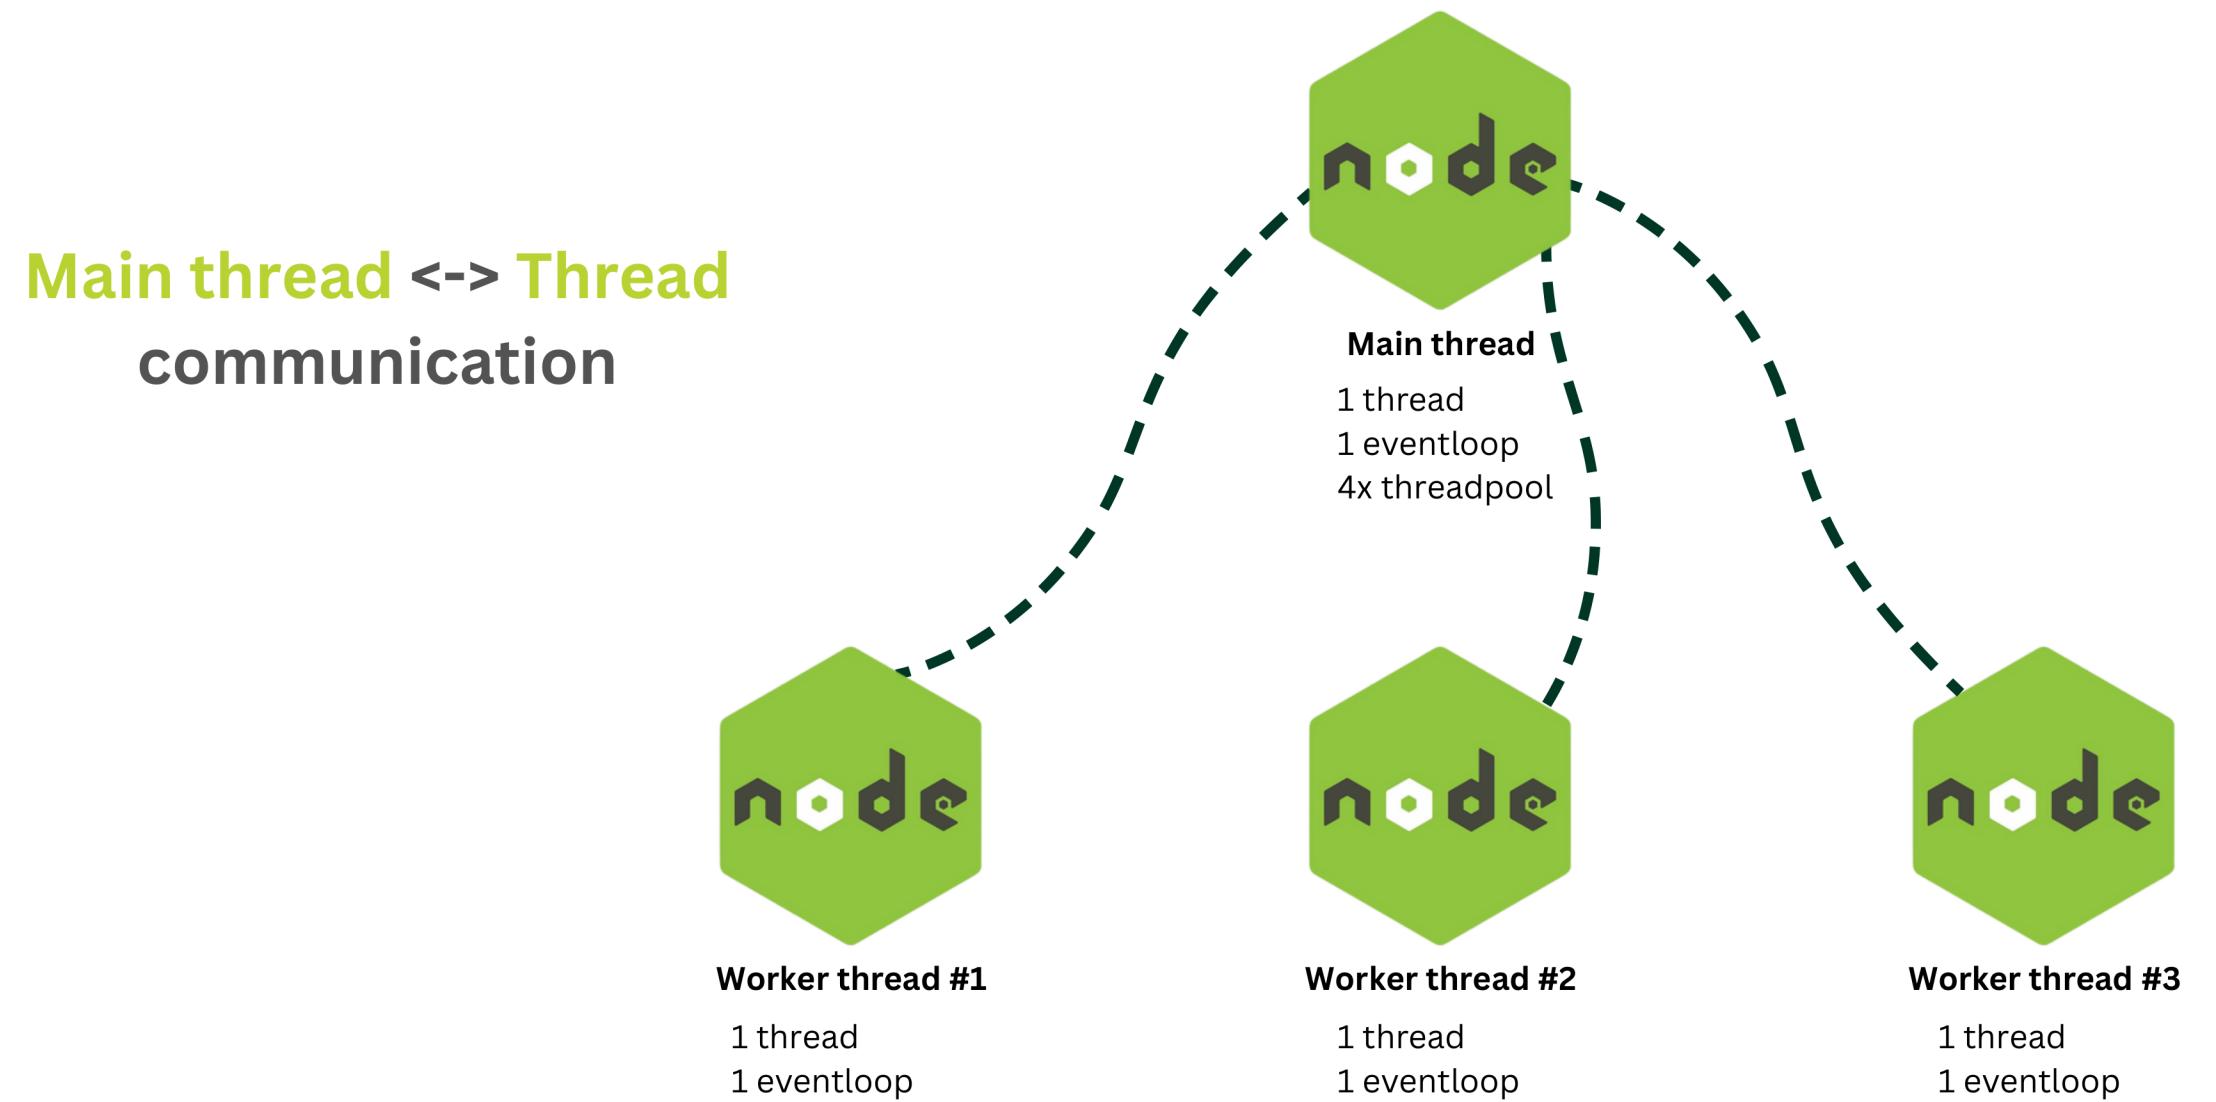
```
const { Worker } = require('worker_threads');
```

Unlike ThreadPool threads...

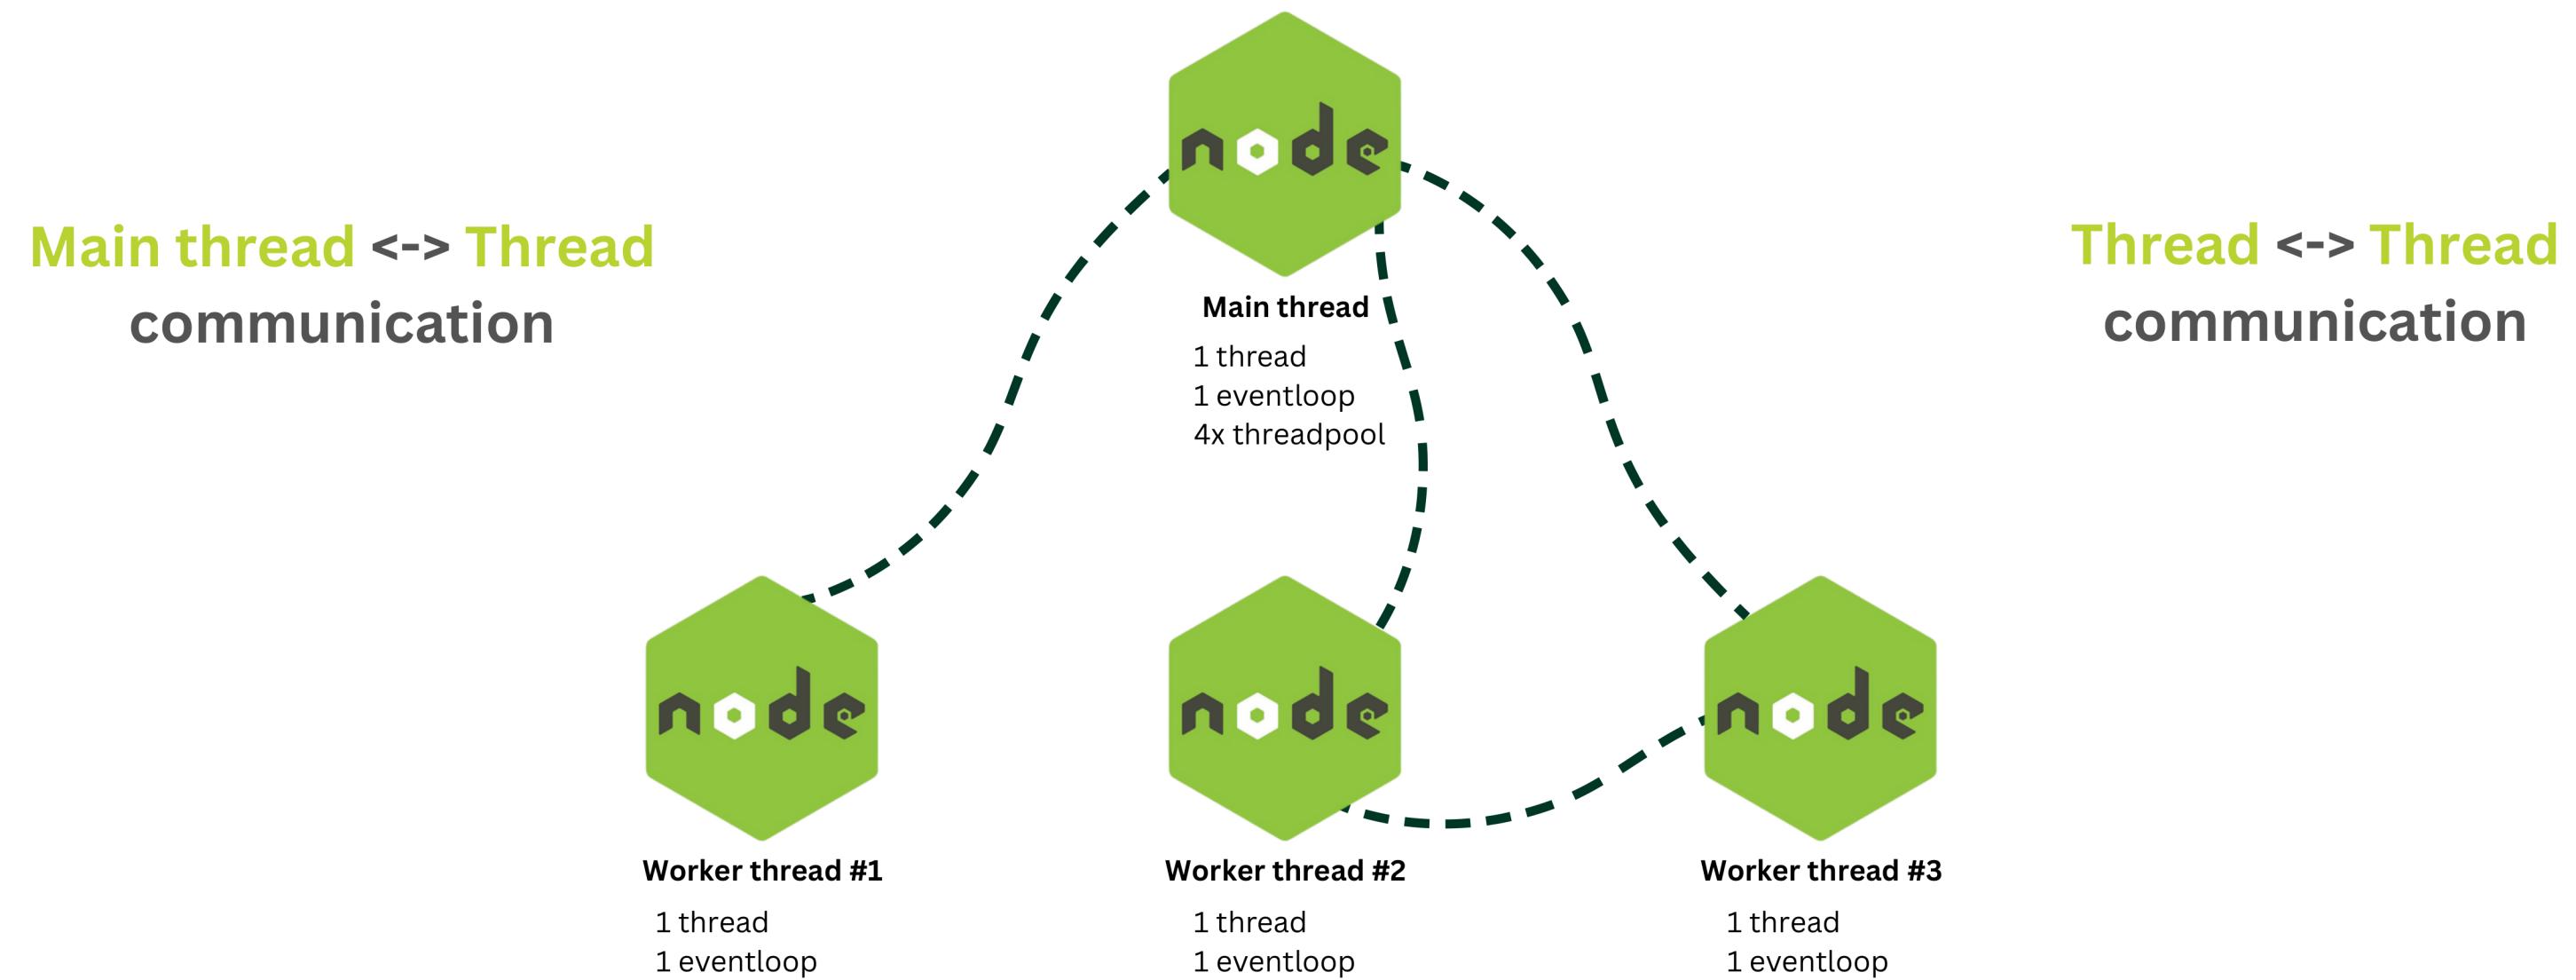
'worker_threads'



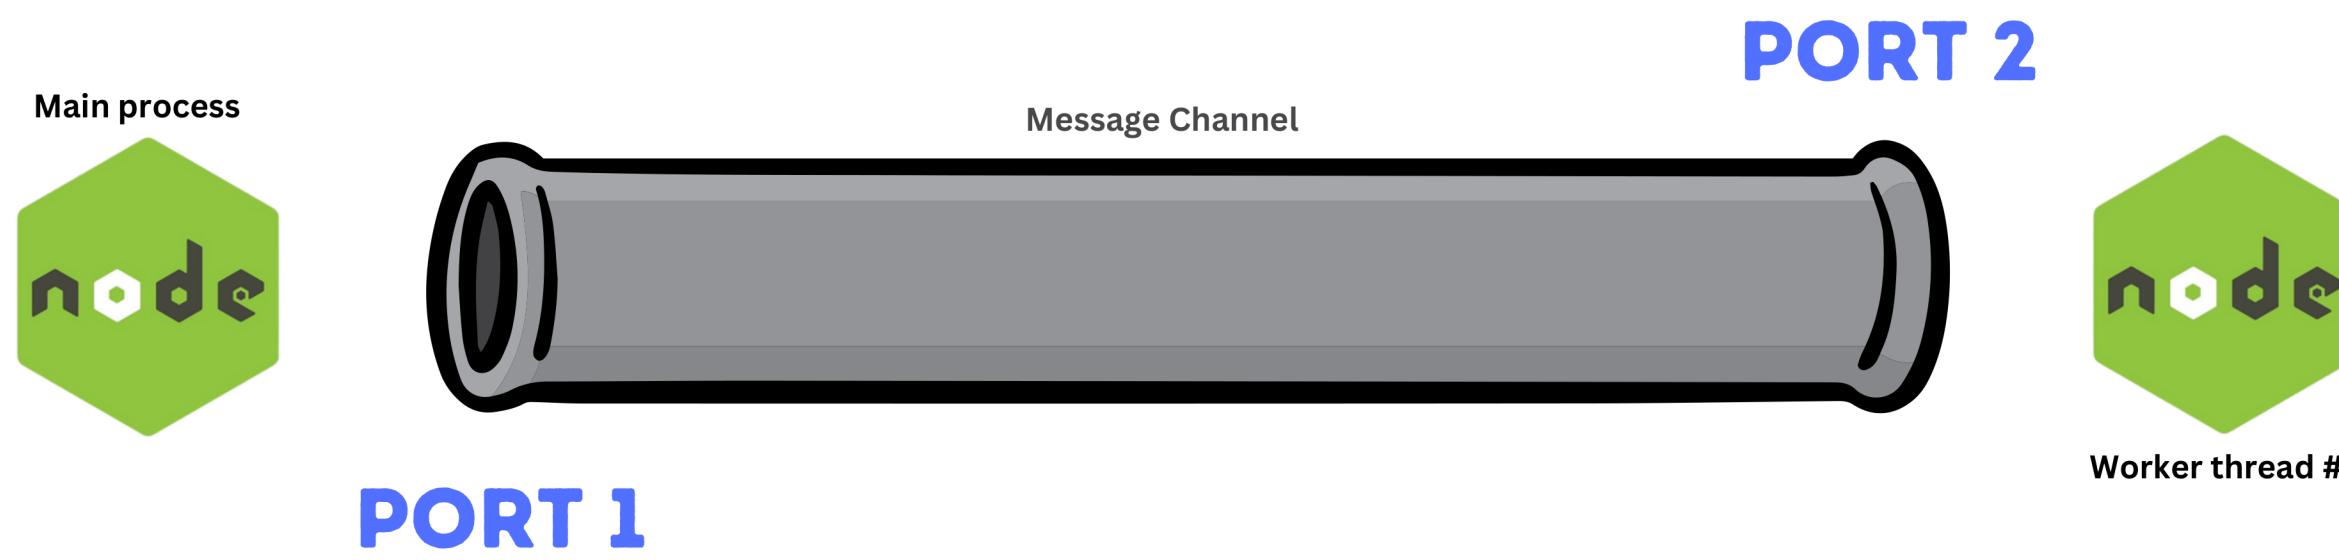
'worker_threads'

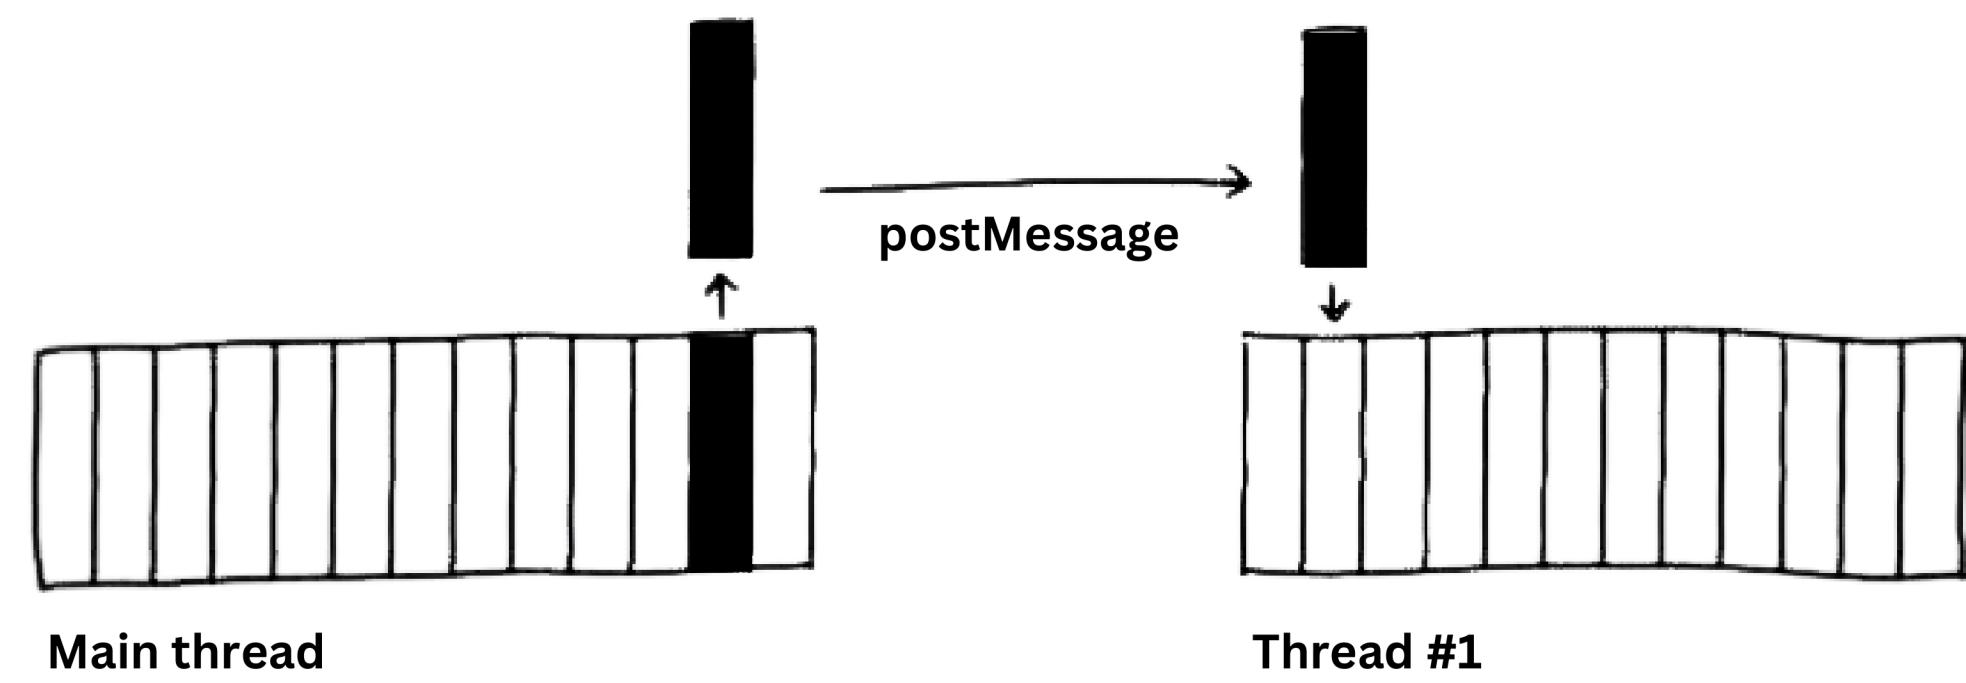


'worker_threads'

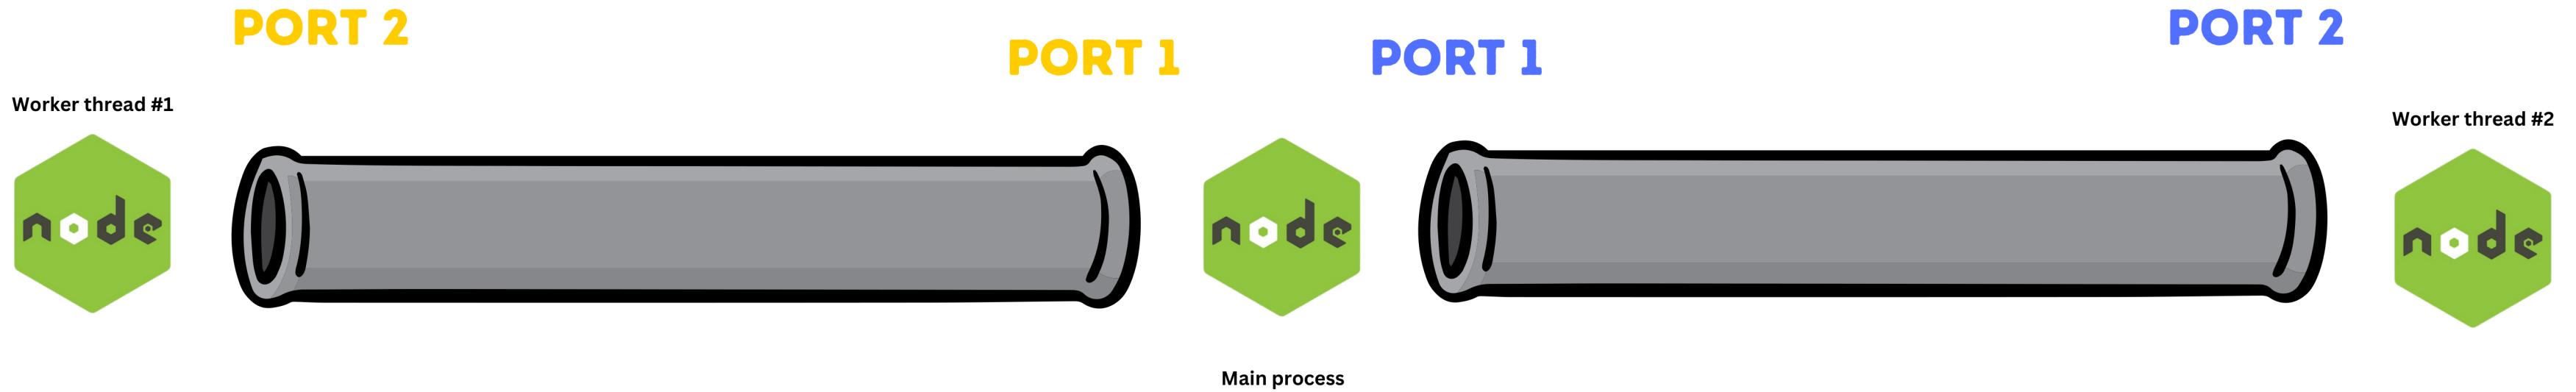


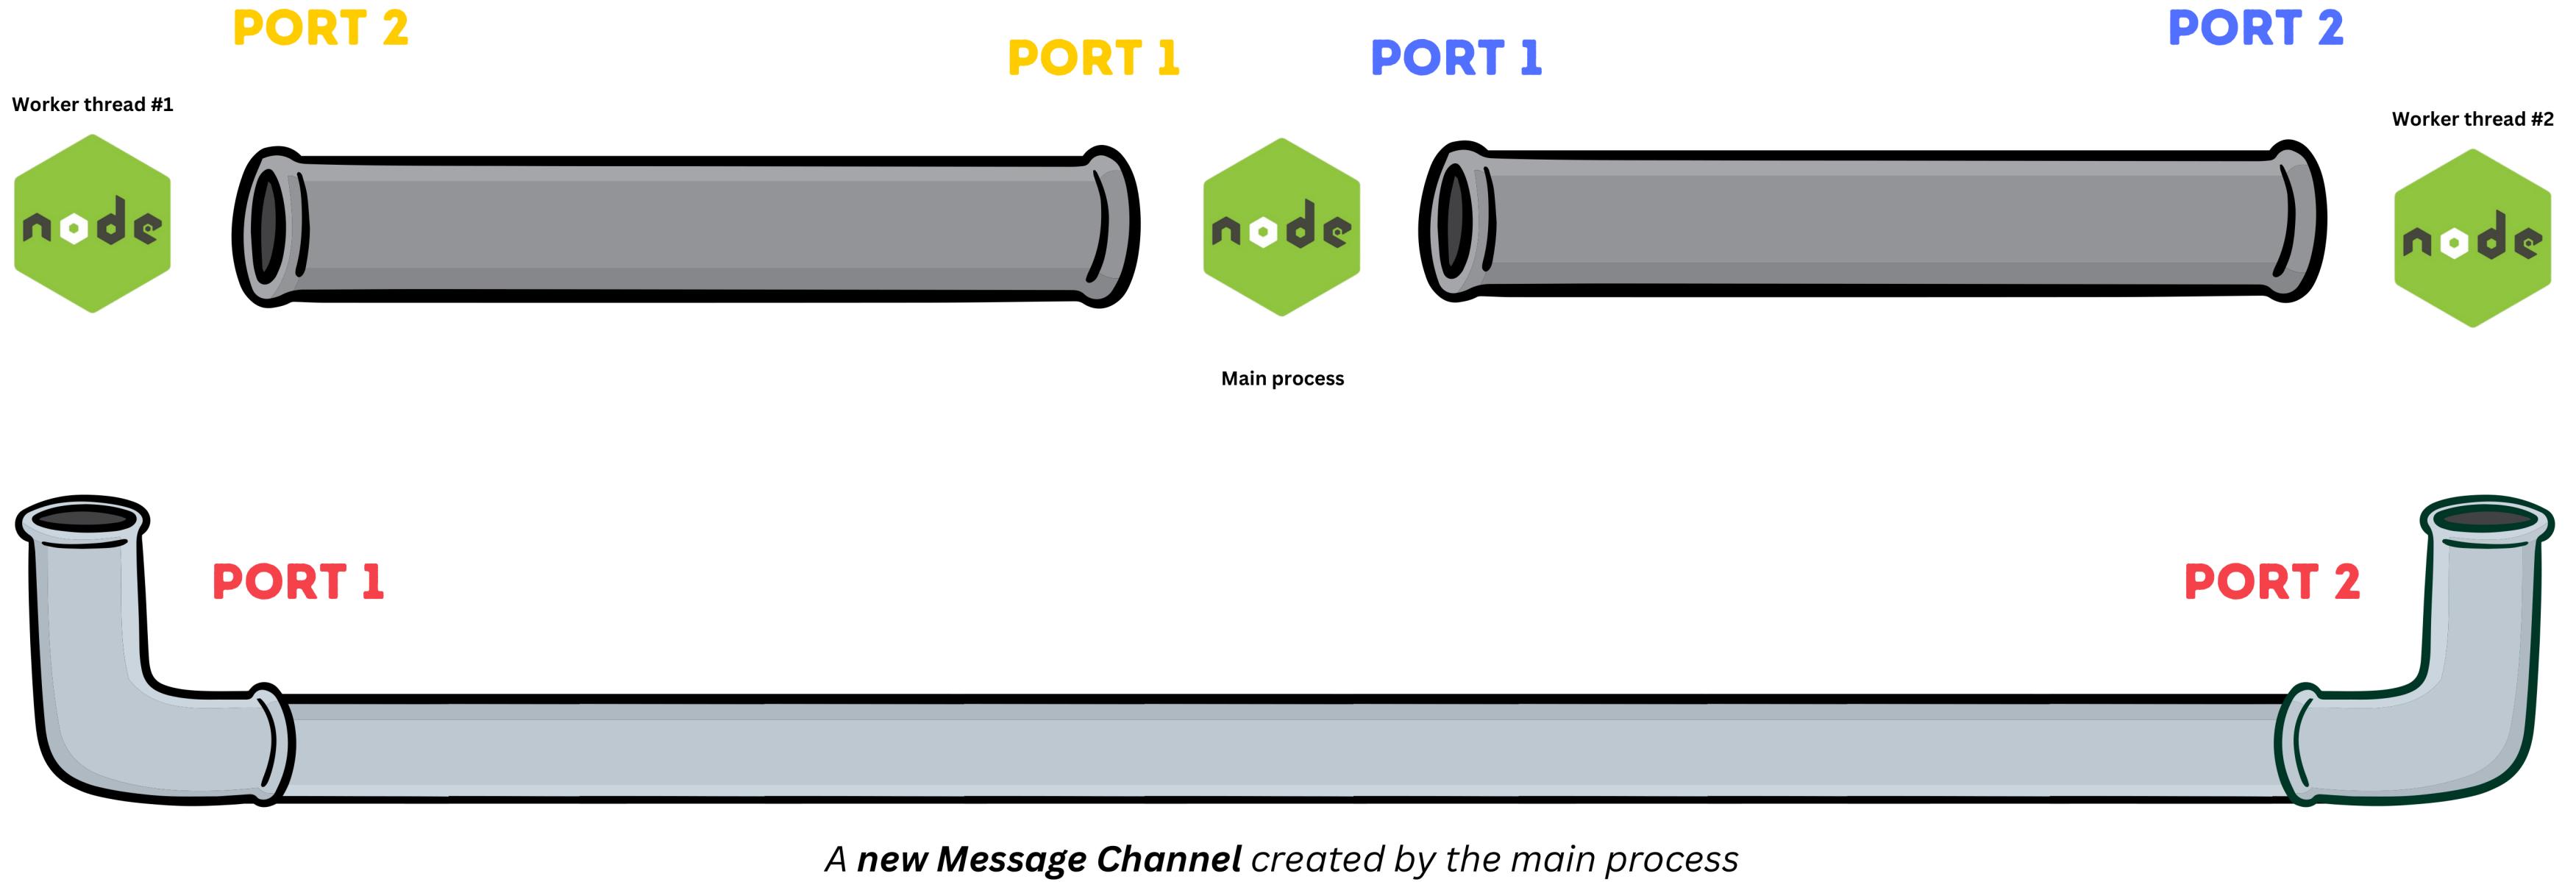
DEMO





Message Channels





Abstractions



Basics

```
// master.js
import { spawn, Thread, Worker } from "threads"

const auth = await spawn(new Worker("./workers/auth"))
const hashed = await auth.hashPassword("Super secret password", "1234")

console.log("Hashed password:", hashed)

await Thread.terminate(auth)
```



```
// workers/auth.js
import sha256 from "js-sha256"
import { expose } from "threads/worker"

expose({
  hashPassword(password, salt) {
    return sha256(password + salt)
  }
})
```



funthreads

A simple library that provides an abstraction for the Node.js `worker_threads` module. It enables you to run your function in a separate thread. You receive a Promise that resolves with the result of your function.

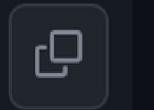
Example

```
const { executeInThread } = require('funthreads');

// heavy operation (this will not block the main thread)
const num = await executeInThread((limit) => {
  let result = 0, i = 1;

  while (i <= limit) {
    result += i.toString().split('').reverse().join('').length;
    i++;
  }

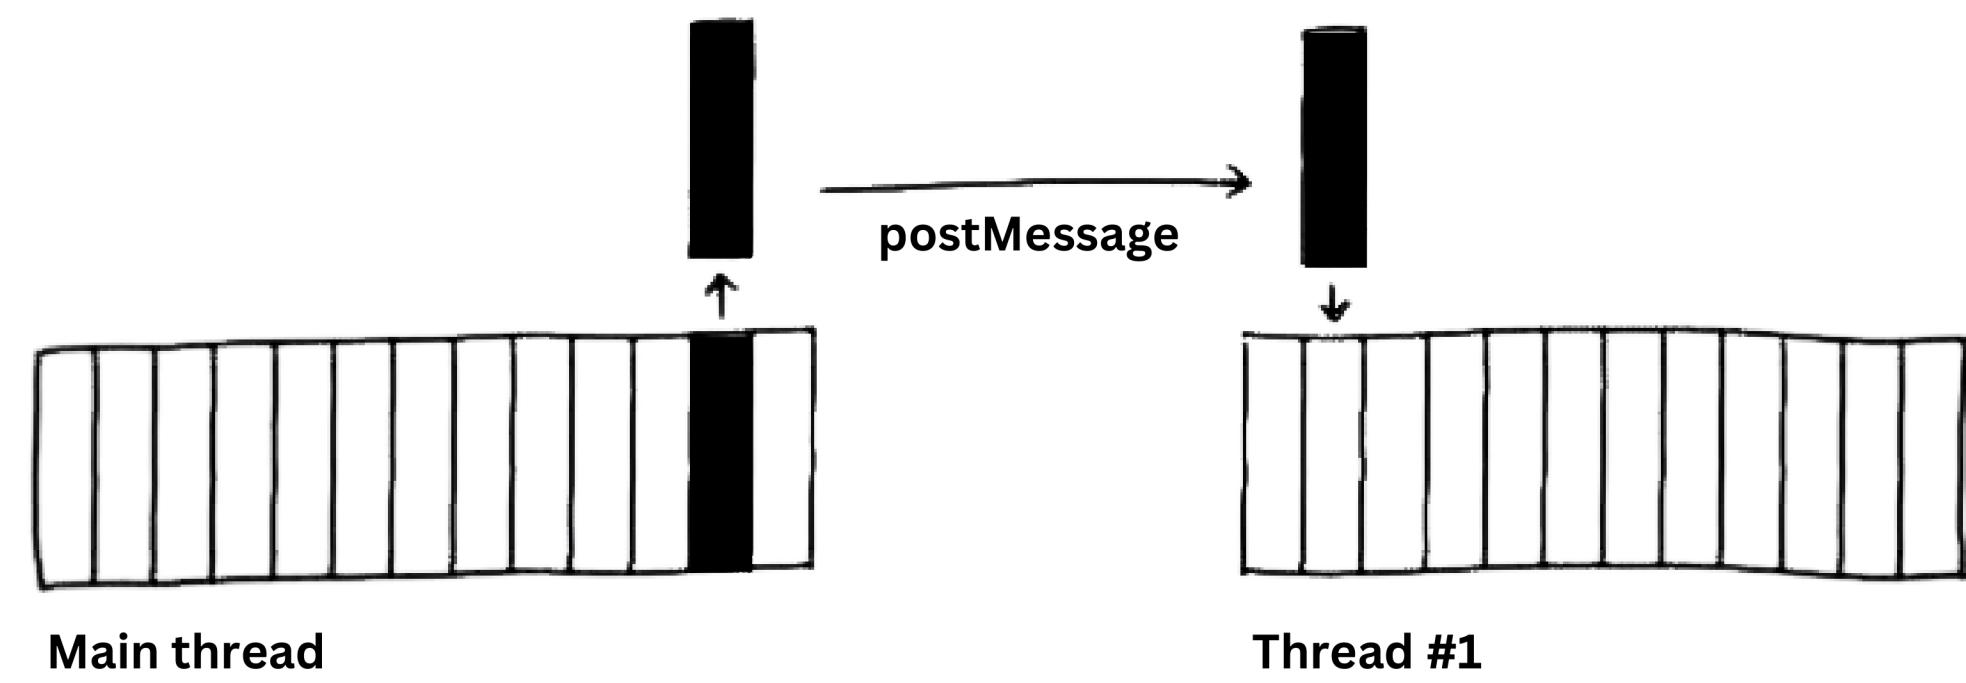
  return result;
}, 12345678);
```



This example highlights the optimization of a resource-intensive calculation. By executing the function in a separate thread, we prevent the main thread from being blocked.

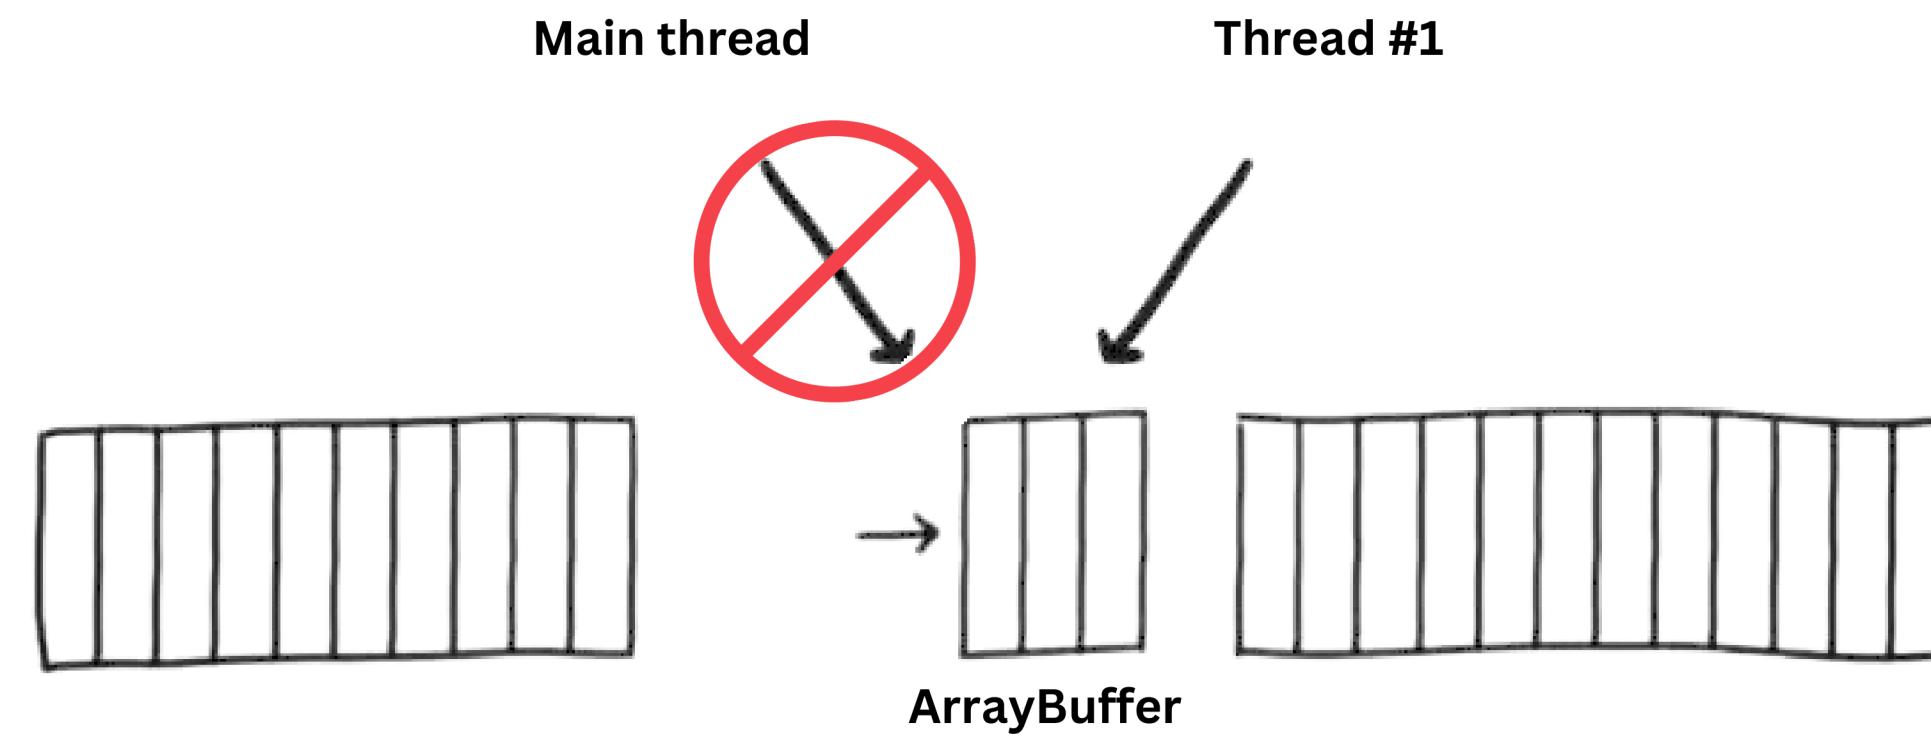
Surprisingly simple, isn't it?

How to share a state/data?



ArrayBuffer

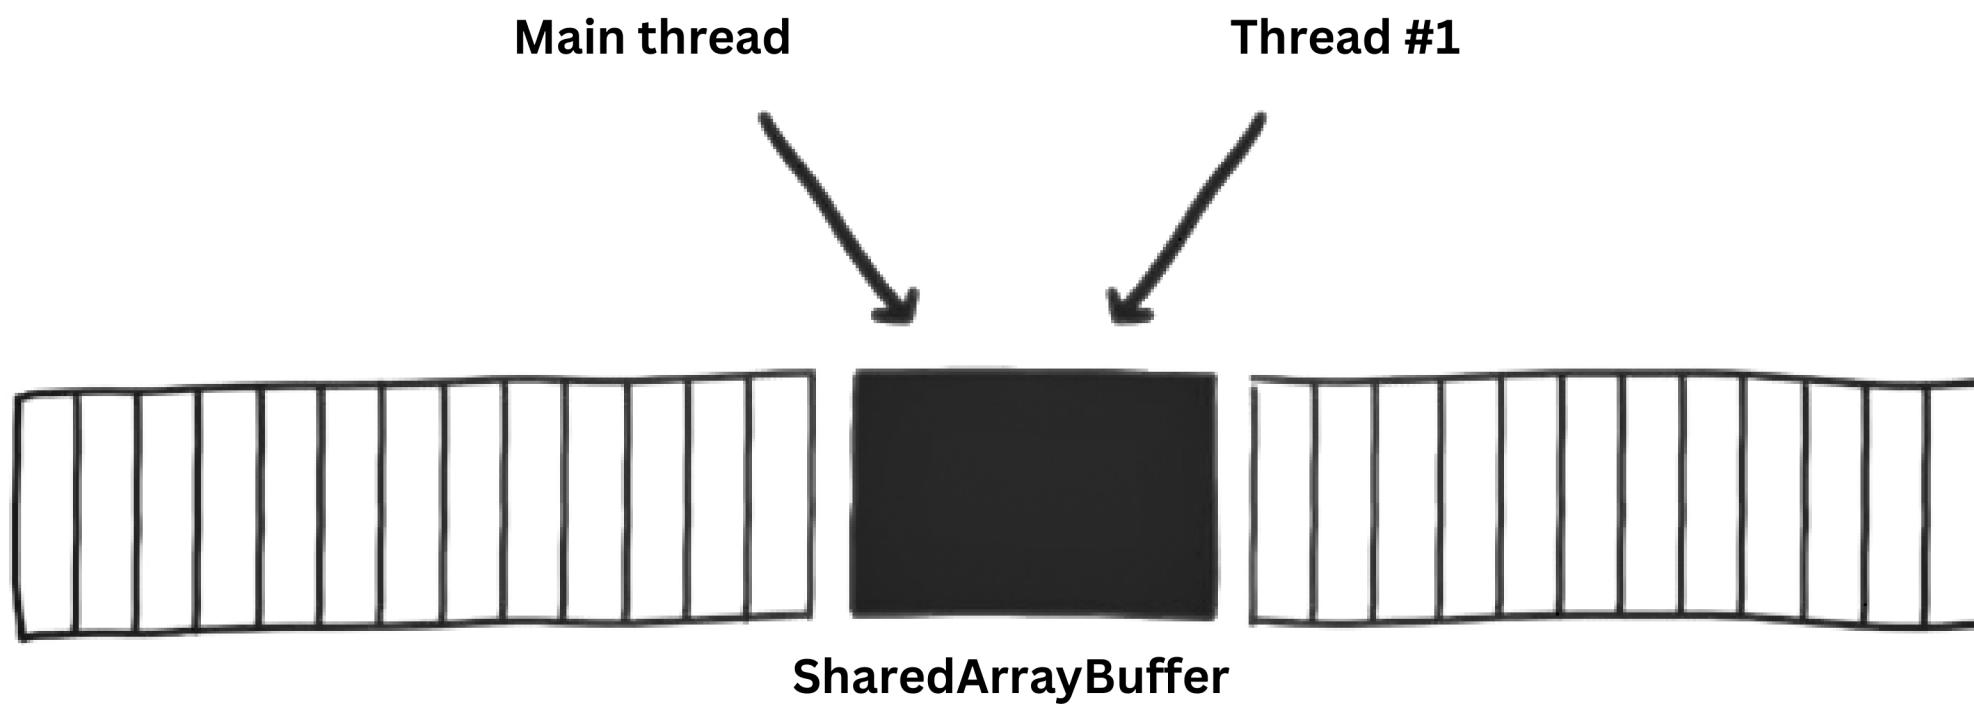
This object is used to represent a generic raw binary data buffer.



SharedArrayBuffer

This object allows thread communication by sharing a common memory space.

It is not a Transferable Object, unlike an ArrayBuffer which is transferable.



01001011010000110...

Typed Arrays(Buffer views)

`Int8Array`

`Uint8Array`

`Uint8ClampedArray`

`Int16Array`

`Uint16Array`

`Int32Array`

`Uint32Array`

`Float32Array`

`Float64Array`

`BigInt64Array`

`BigUint64Array`

Int8Array

The **Int8Array** typed array represents an array of **8-bit** signed integers.

It can store values ranging from **-128 to 127**, as it represents 8-bit signed integers.

Uint8Array

An array of **8-bit unsigned integers**, meaning it can store values from **0 to 255**.

Int16Array

Uint16Array

Int32Array

Uint32Array

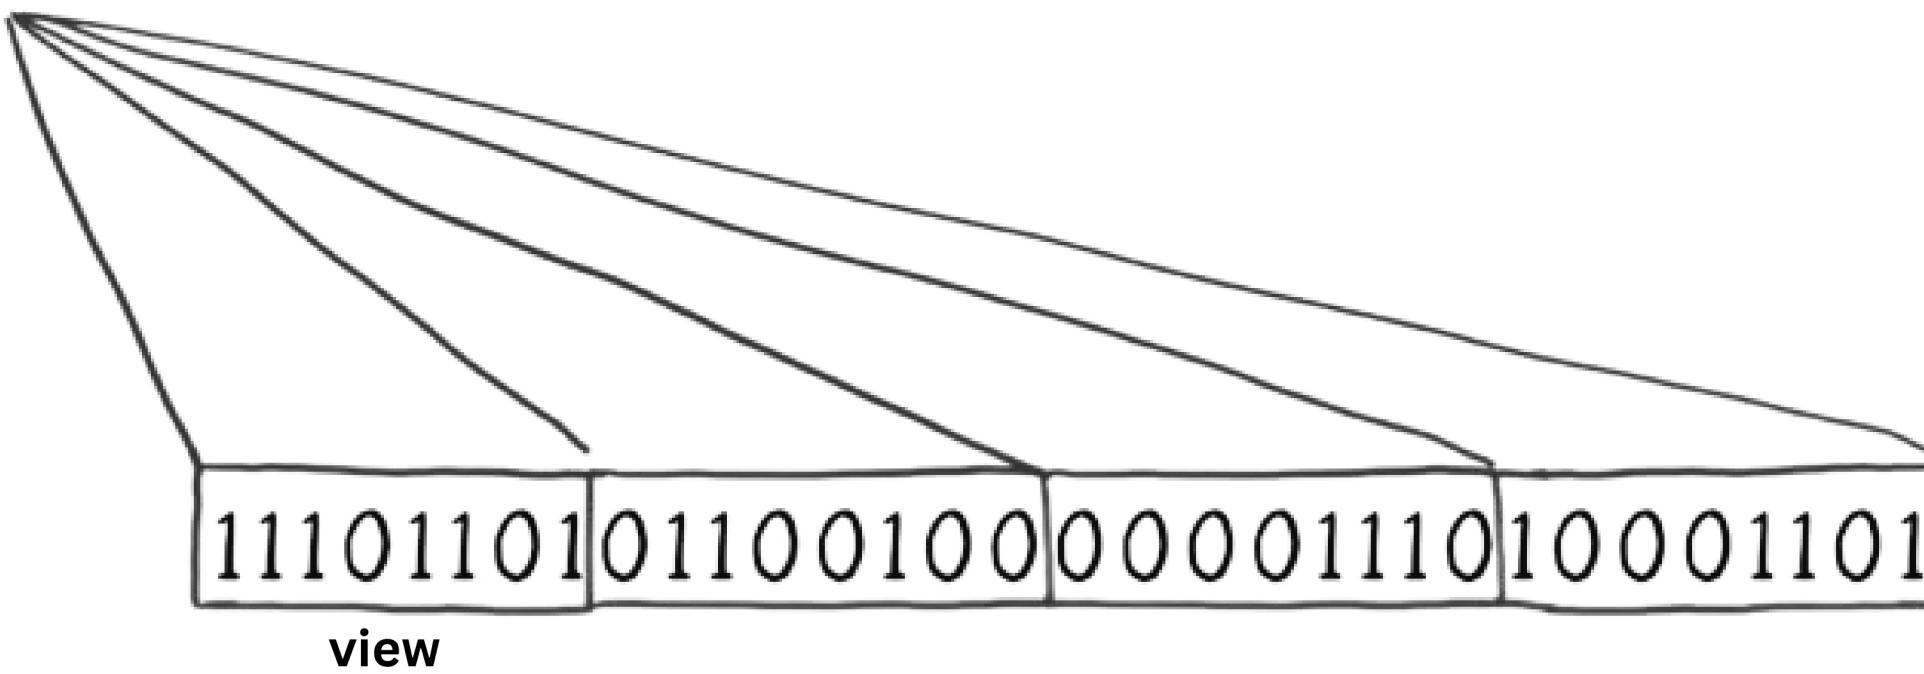
Float32Array

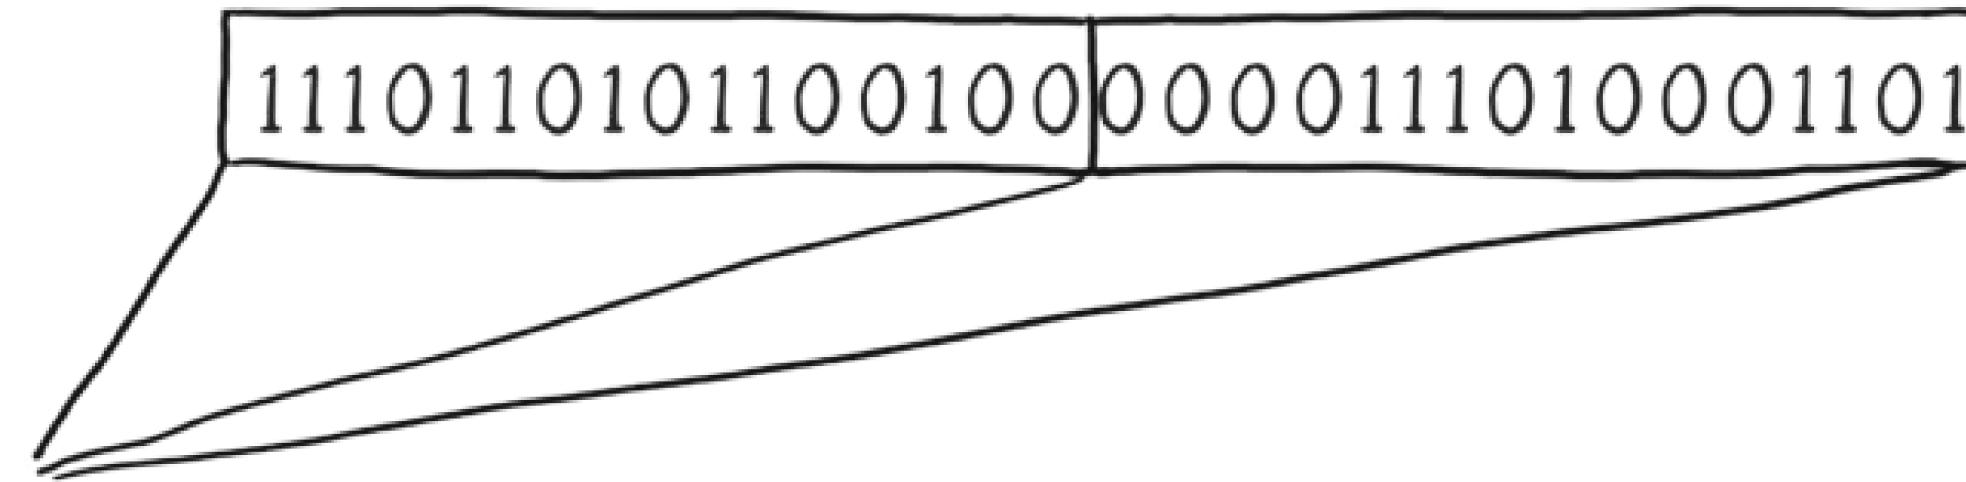
Float64Array

BigInt64Array

BigUint64Array

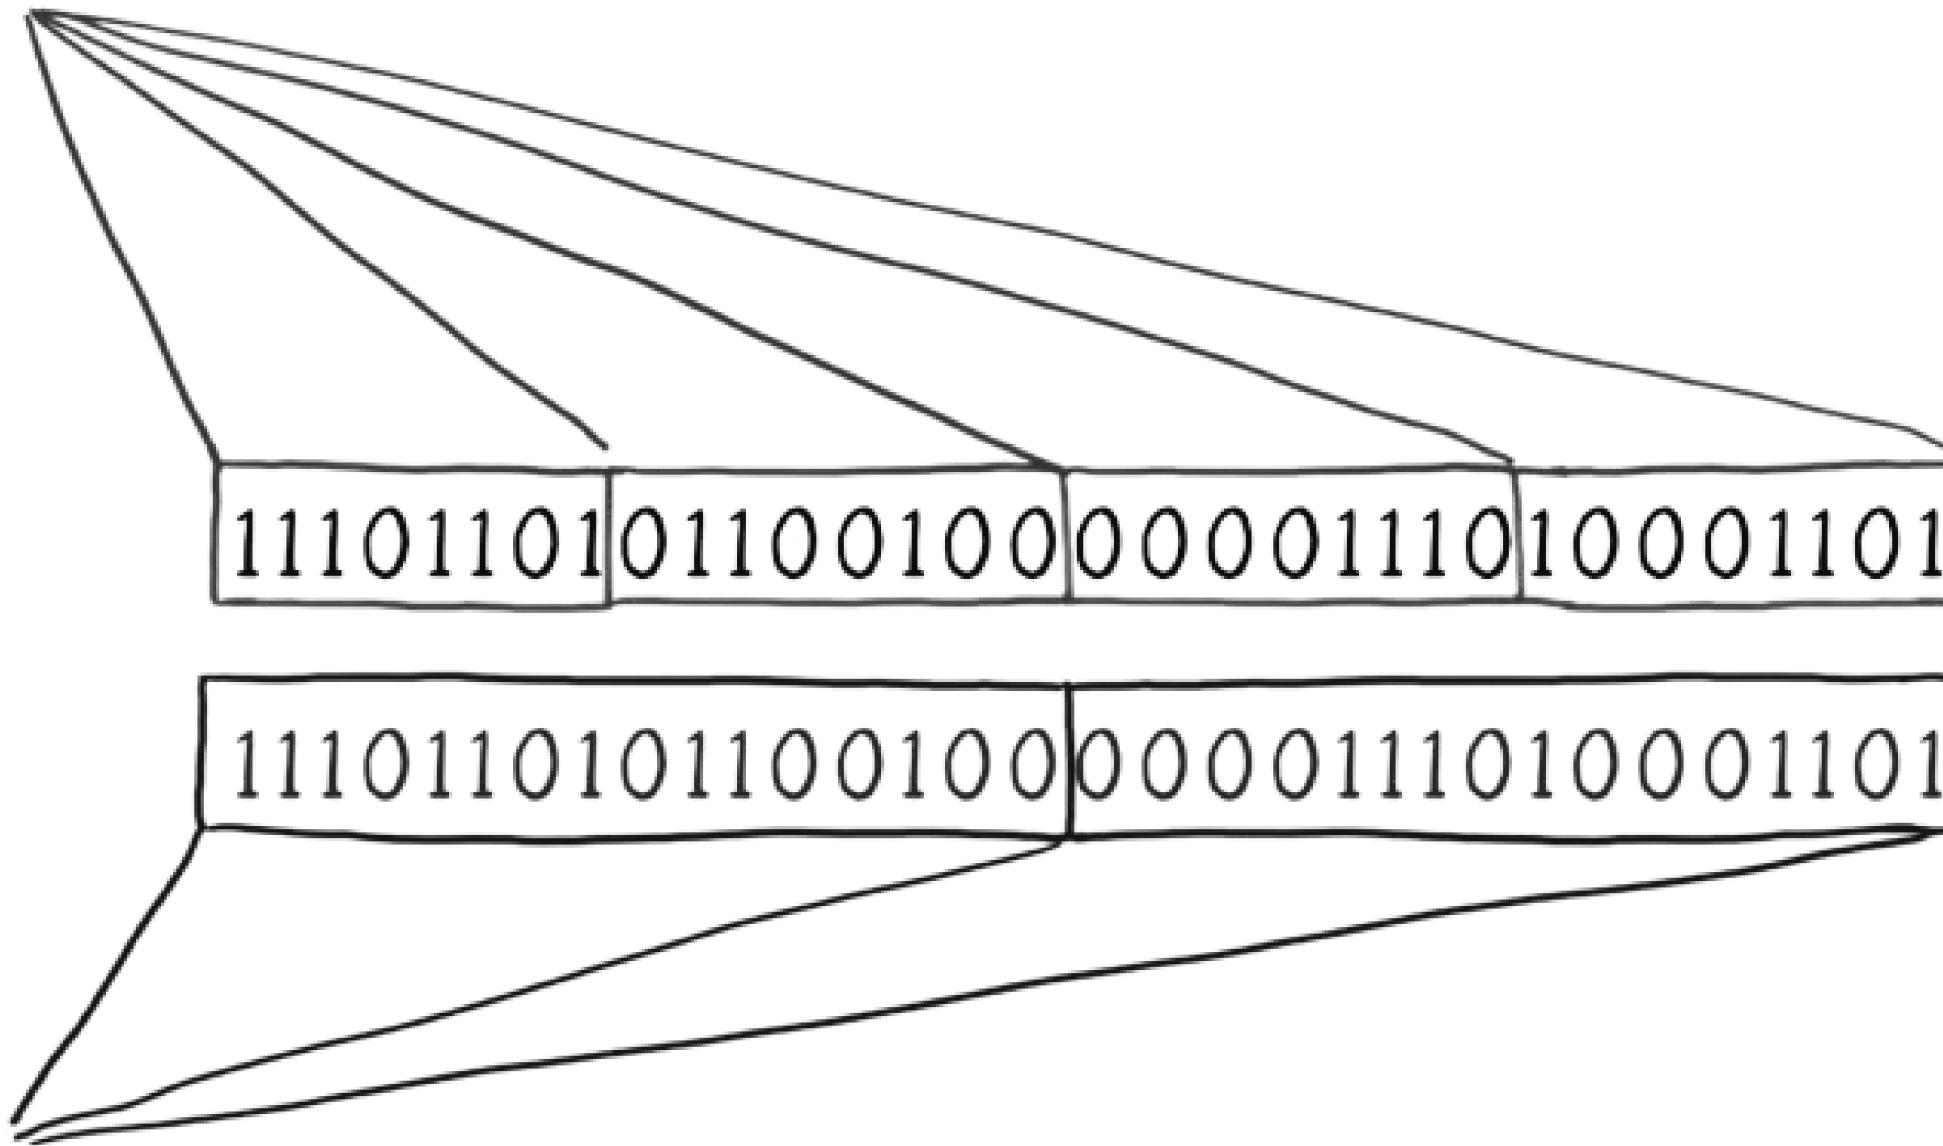
Int8Array





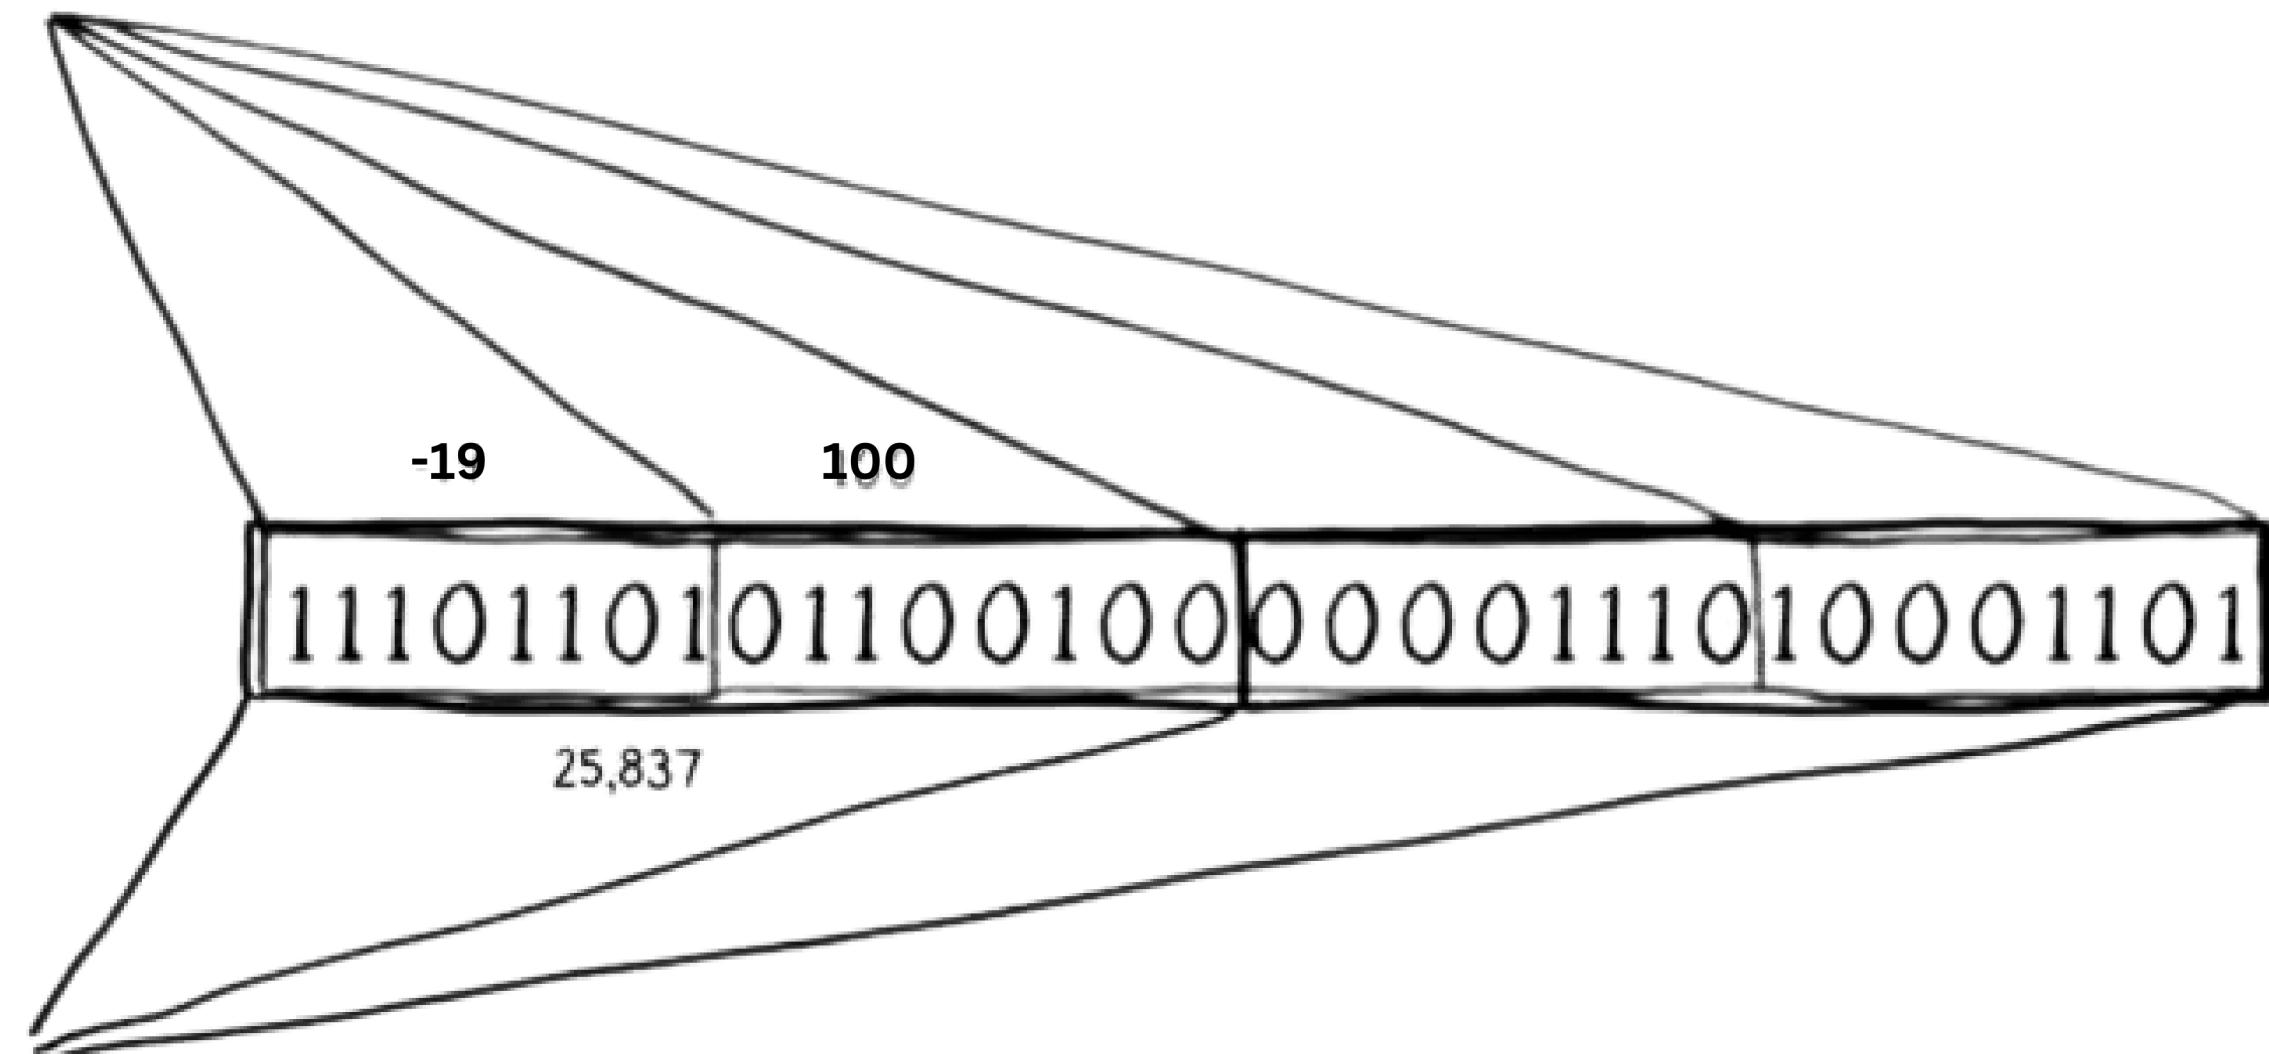
`Uint16Array`

Int8Array



Uint16Array

Int8Array



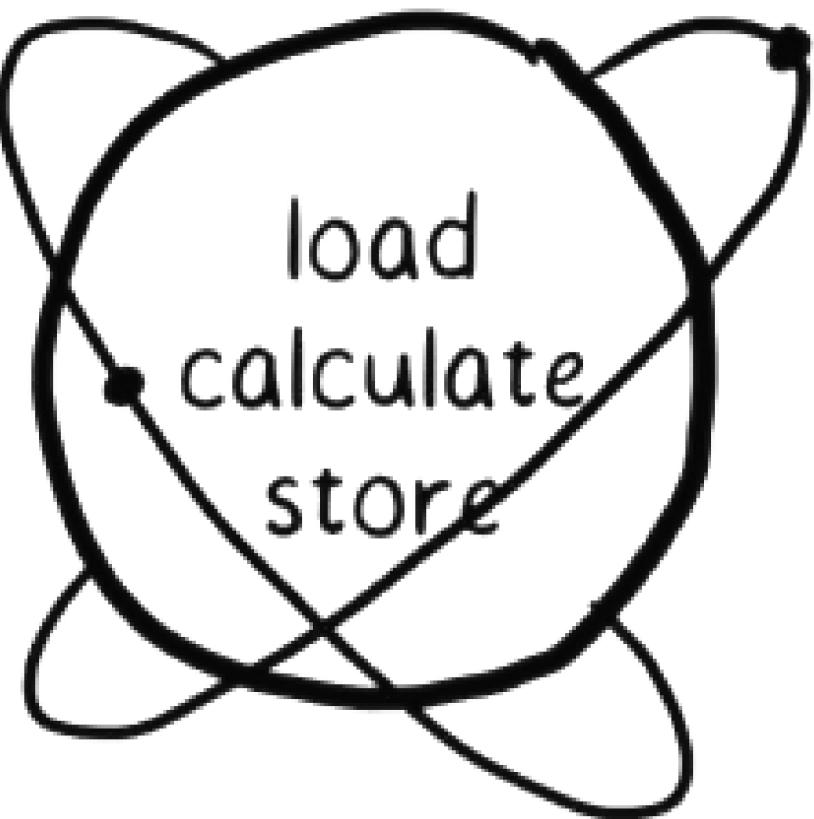
How to keep data using numbers?

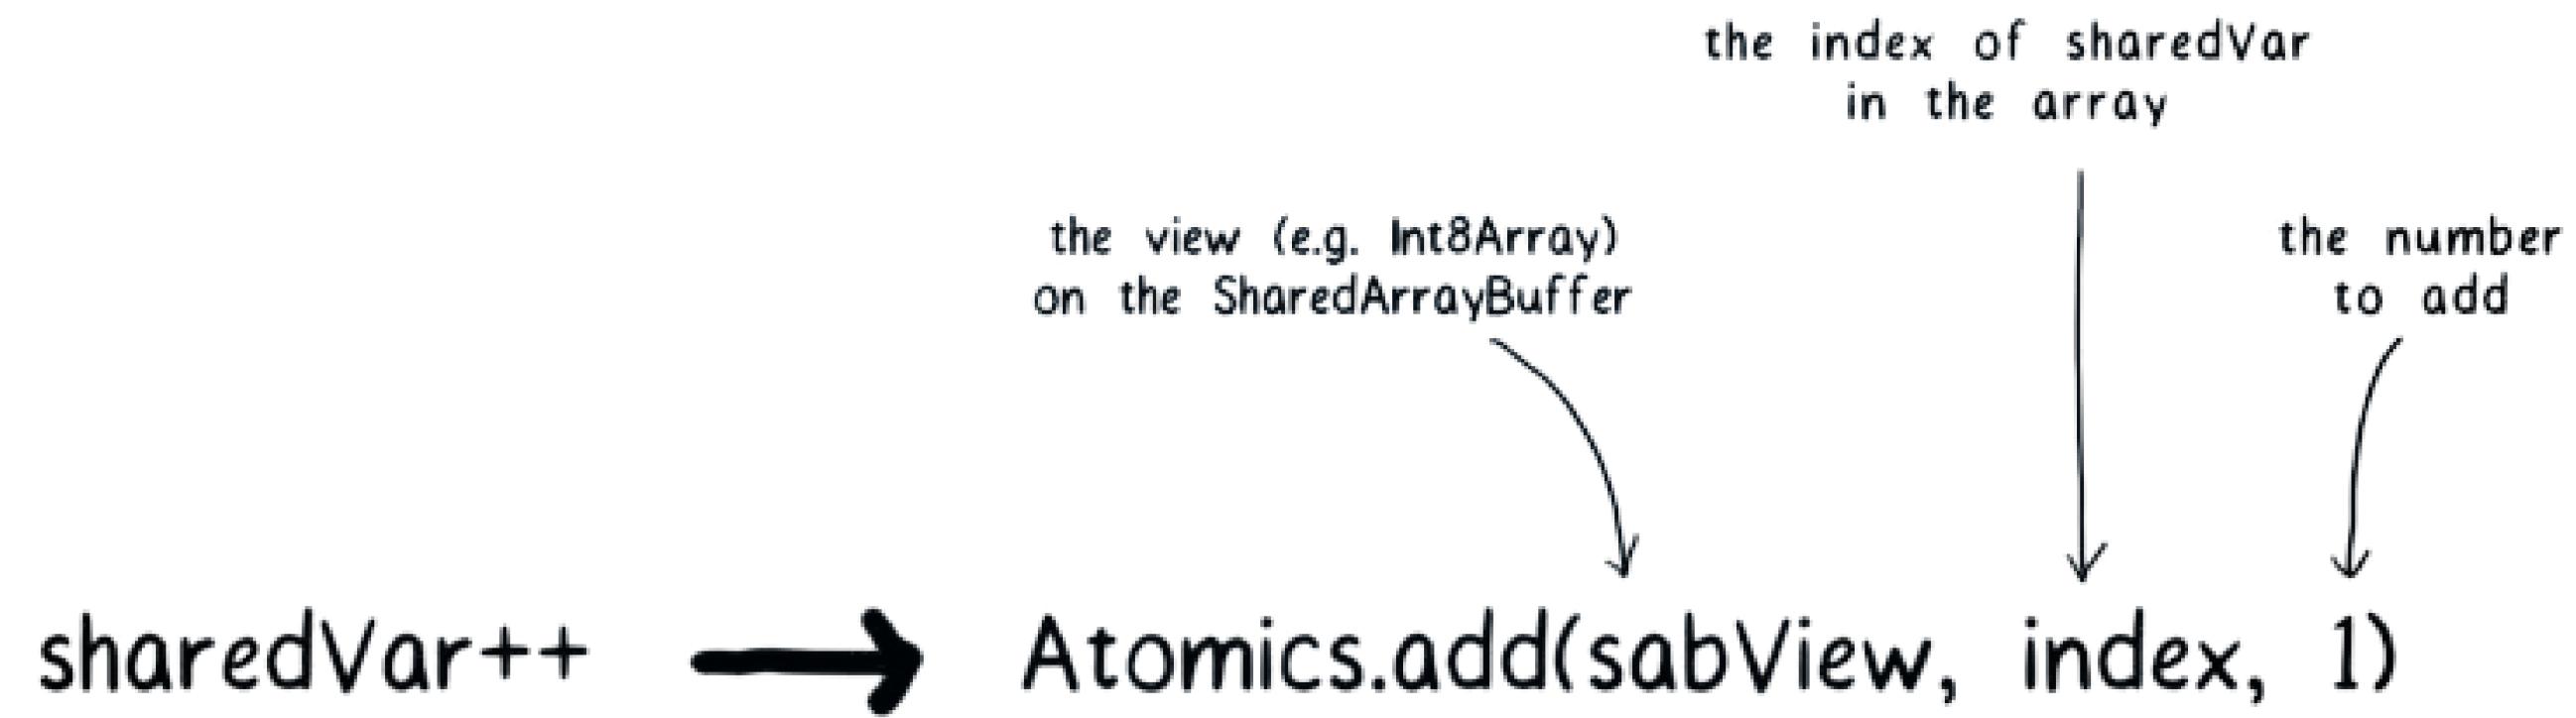
Race conditions

Atomics

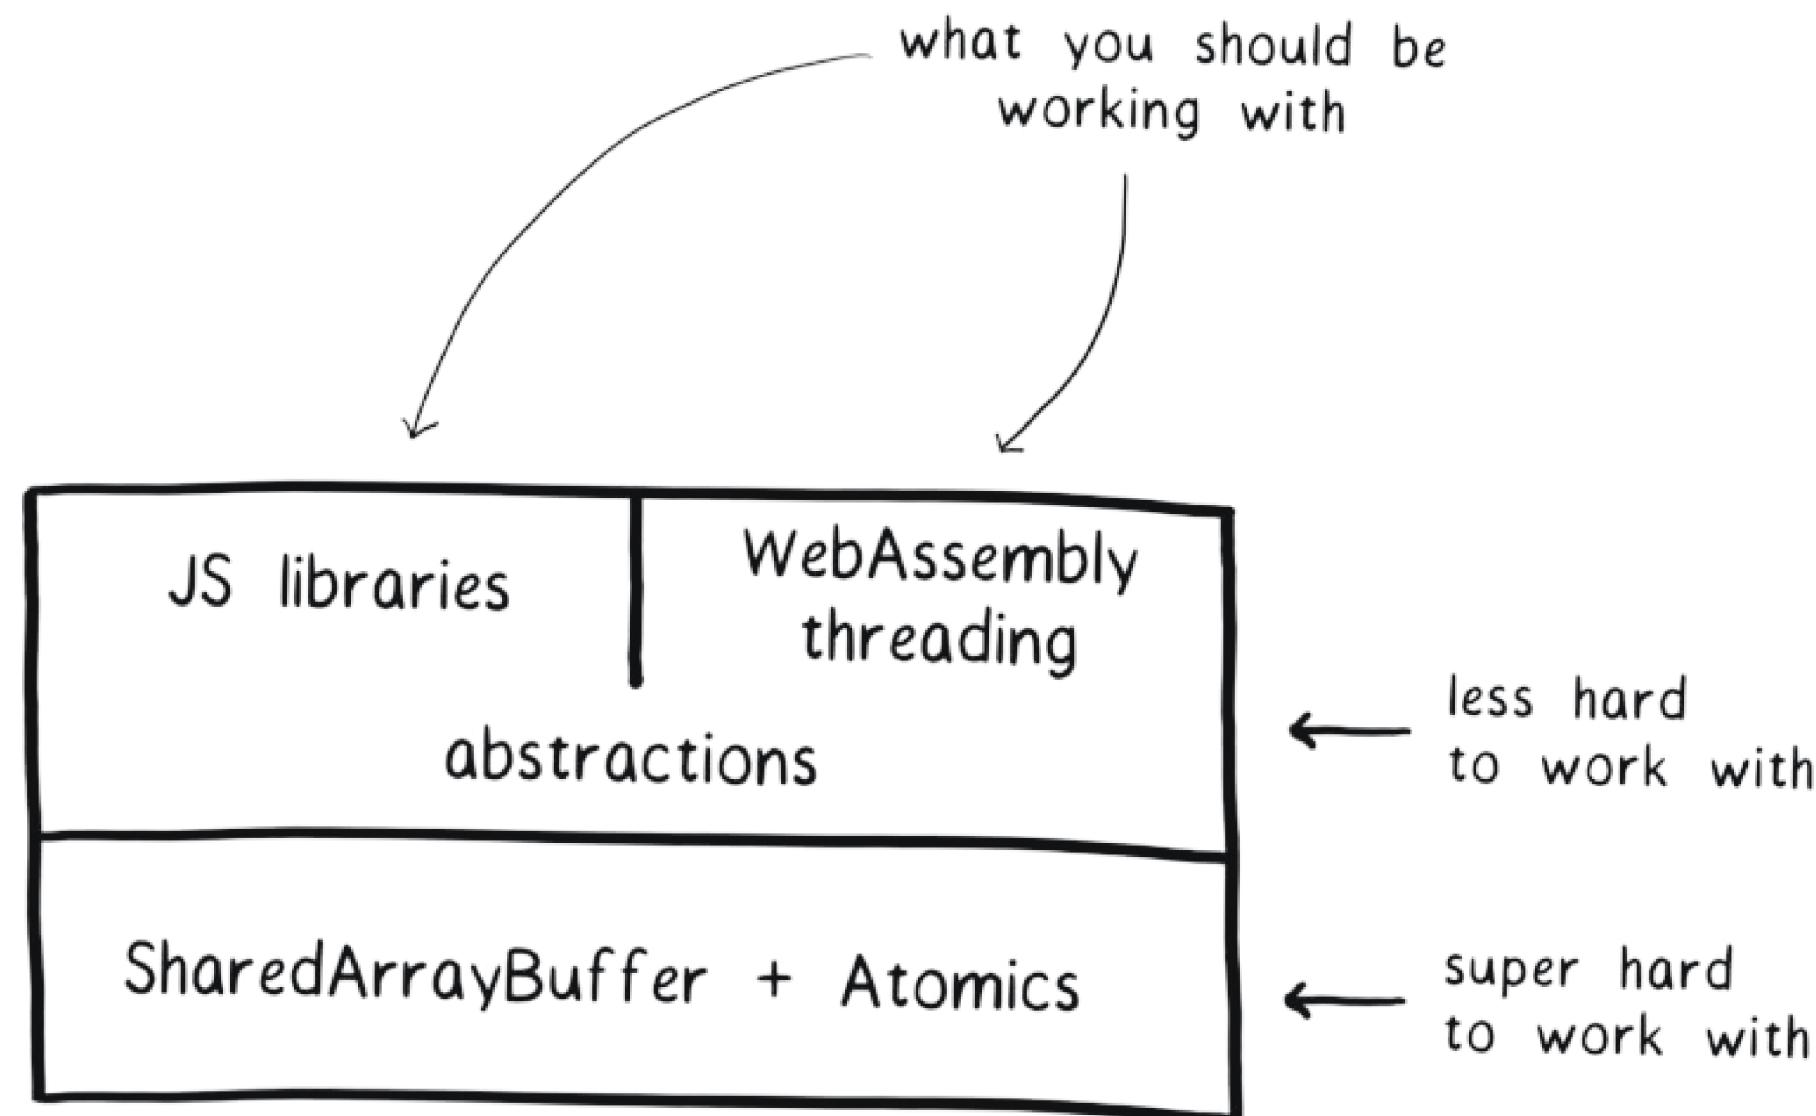
This object contains static methods for carrying out atomic operations.

They are used with SharedArrayBuffer and ArrayBuffer objects.





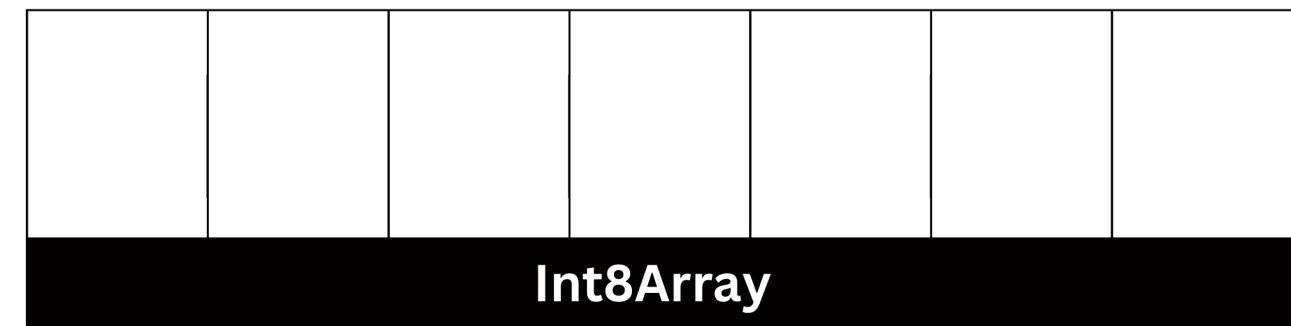
Atomics and SharedBuffers can be pretty challenging to handle.



How to keep data using numbers?

```
const b = new SharedArrayBuffer( 7 );  
const v = new Int8Array( b );
```

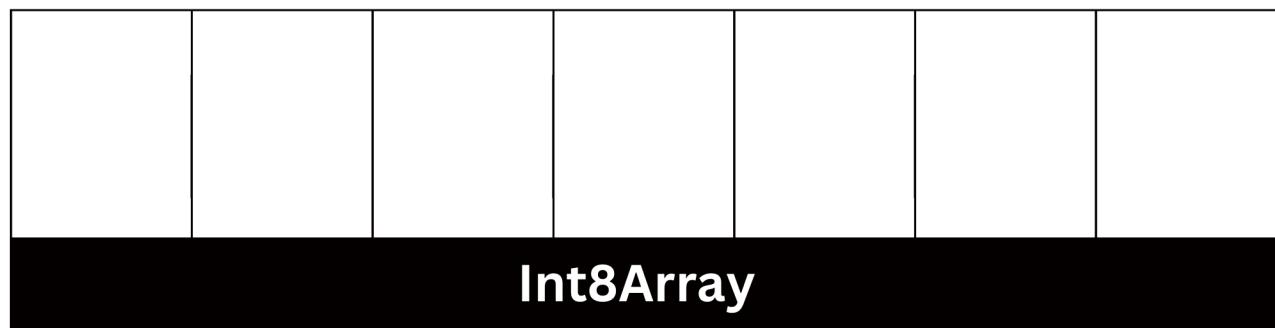
SharedArrayBuffer



```
const b = new SharedArrayBuffer( 7 );  
const v = new Int8Array( b );
```

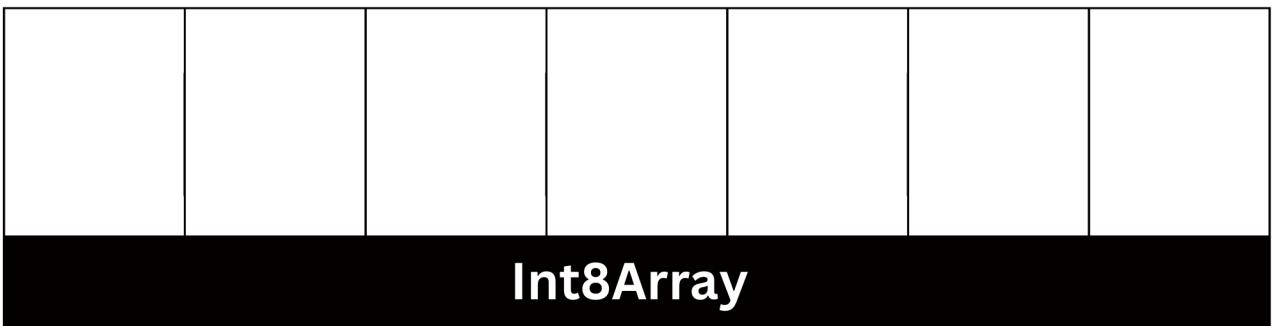
h e l l o

SharedArrayBuffer



```
const b = new SharedArrayBuffer( 7 );  
const v = new Int8Array( b );
```

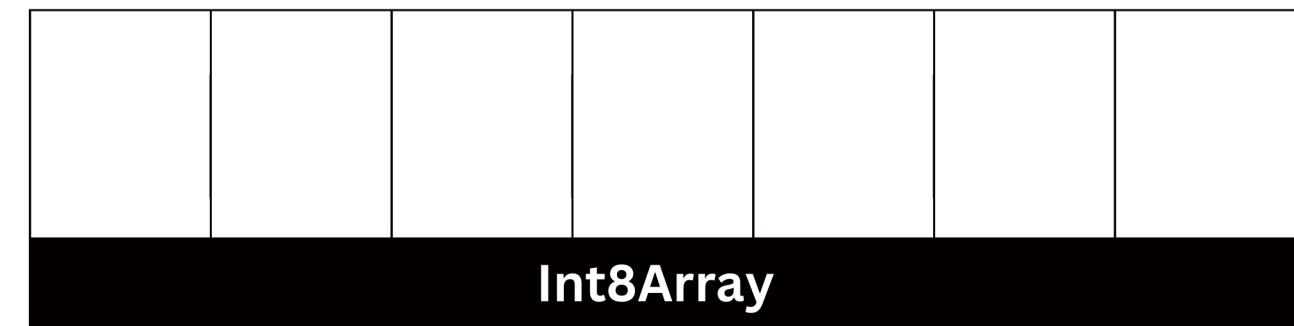
SharedArrayBuffer



h e l l o
104 101 108 108 111

```
const b = new SharedArrayBuffer( 7 );  
const v = new Int8Array( b );
```

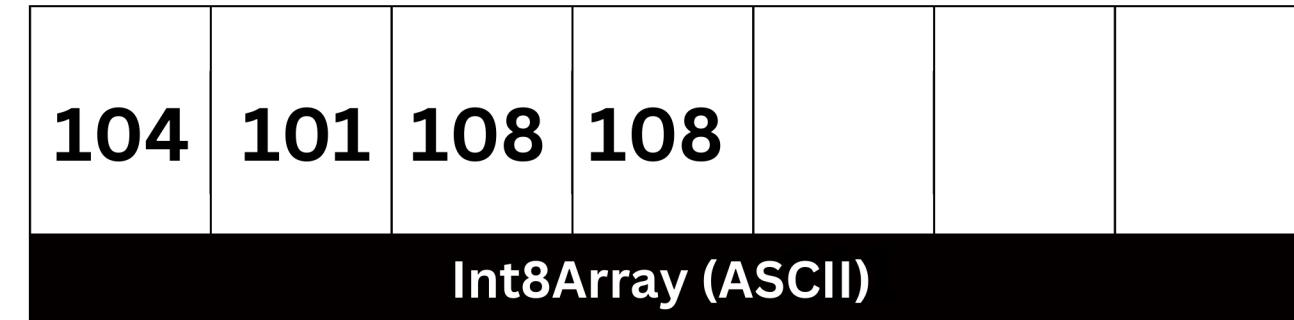
SharedArrayBuffer



h e l l o
104 101 108 108 111

```
v[0] = 104;  
v[1] = 101;  
...  
...
```

SharedArrayBuffer



First, let's create a simple hash map structure in main process, then create a worker thread and share the hash.

```
// main.js
const { Worker } = require('worker_threads');
const { WorkerMap } = require('worker_map');

const map = new WorkerMap();
map.set('balance', 100); // sync operation

new Worker('./worker.js', {
  workerData: {
    mapBuffer: map.toSharedBuffer(),
  },
});

setTimeout(() => {
  console.log(map.get('balance')); // 200
}, 50);
```



Now, let's access the shared hash map structure in the worker thread.

```
// worker.js
const { WorkerMap } = require('worker_map');
const { workerData } = require('worker_threads');

const map = new WorkerMap(workerData.mapBuffer);
console.log(map.get('balance')); // 100

// The change will be reflected in the main process as well
map.set('balance', 200);
```



worker_map ↵

Instance methods ↵

Worker_map is much like JavaScript's regular Map.

map.set(key, value) ↵

```
map.set('name', 'John'); // true
```

map.get(key): ↵

```
const name = map.get('name'); // 'John'
```

map.delete(key): ↵

```
map.delete('name'); // true  
map.delete('something'); // false because it doesn't exist
```

map.clear(): ↵

```
map.clear();  
map.size(); // 0
```

map.has(key) ↵

```
map.has('name'); // true  
map.has('country'); // false
```

map.size() ↵

```
map.has('size'); // 1
```

map.keys() ↵

```
map.keys(); // [ 'name' ]
```

map.entries() ↵

```
for (const [ key, value ] of map.entries()) {  
  console.log(`key: ${key}`); // name: 'John'  
}
```

map.forEach() ↵

```
map.forEach(function(key, value, map) {  
  console.log(`key: ${key}`); // name: 'John'  
});
```

map.toSharedBuffer() ↵

```
const buffer = map.toSharedBuffer();  
const sameMap = new WorkerMap(buffer);
```

map.toObject() ↵

```
const mapObject = map.toObject(); // { ... }  
mapObject.name; // 'John'
```

worker_map ↵

Instance methods ↵

Worker_map is much like JavaScript's regular Map.

map.set(key, value) ↵

```
map.set('name', 'John'); // true
```

map.get(key): ↵

```
const name = map.get('name'); // 'John'
```

map.delete(key): ↵

```
map.delete('name'); // true  
map.delete('something'); // false because it doesn't exist
```

map.clear(): ↵

```
map.clear();  
map.size(); // 0
```

map.has(key) ↵

```
map.has('name'); // true  
map.has('country'); // false
```

map.size() ↵

```
map.has('size'); // 1
```

map.keys() ↵

```
map.keys(); // [ 'name' ]
```

map.entries() ↵

```
for (const [ key, value ] of map.entries()) {  
  console.log(`${key}: ${value}`); // name: 'John'  
}
```

map.forEach() ↵

```
map.forEach(function(key, value, map) {  
  console.log(`${key}: ${value}`); // name: 'John'  
});
```

map.toSharedBuffer() ↵

```
const buffer = map.toSharedBuffer();  
const sameMap = new WorkerMap(buffer);
```

map.toObject() ↵

```
const mapObject = map.toObject(); // { ... }  
mapObject.name; // 'John'
```

First, let's create a simple hash map structure in main process, then create a worker thread and share the hash.

```
// main.js
const { Worker } = require('worker_threads');
const { WorkerMap } = require('worker_map');

const map = new WorkerMap();
map.set('balance', 100); // sync operation

new Worker('./worker.js', {
  workerData: {
    mapBuffer: map.toSharedBuffer(),
  },
});

setTimeout(() => {
  console.log(map.get('balance')); // 200
}, 50);
```



Now, let's access the shared hash map structure in the worker thread.

```
// worker.js
const { WorkerMap } = require('worker_map');
const { workerData } = require('worker_threads');

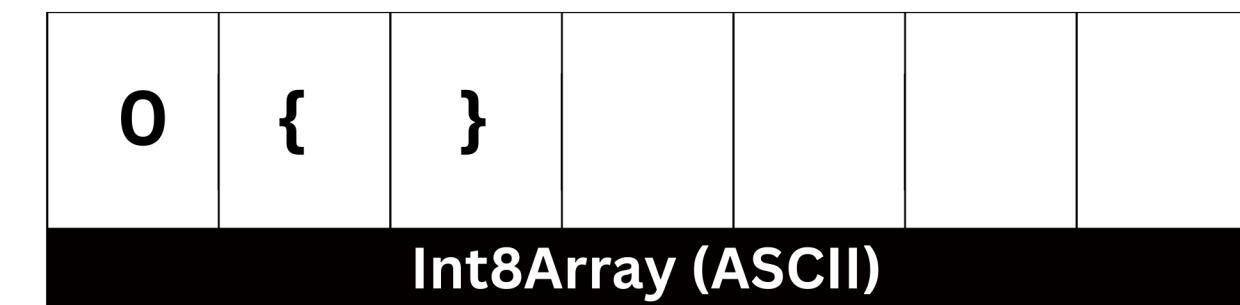
const map = new WorkerMap(workerData.mapBuffer);
console.log(map.get('balance')); // 100

// The change will be reflected in the main process as well
map.set('balance', 200);
```



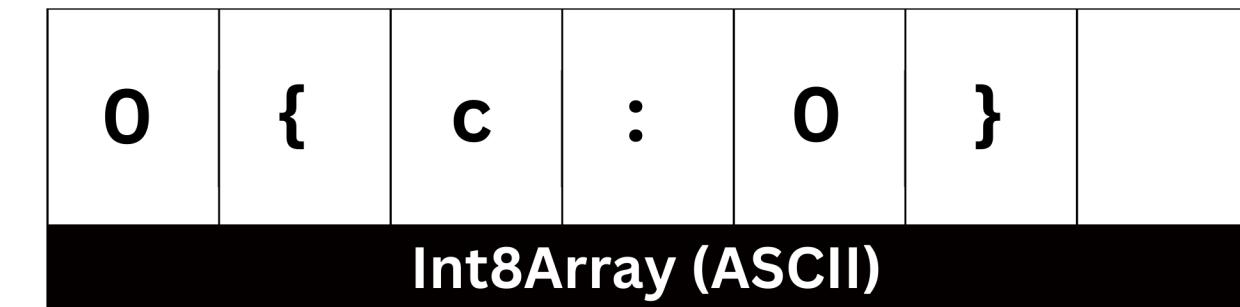
```
const map = new WorkerMap();
```

SharedArrayBuffer



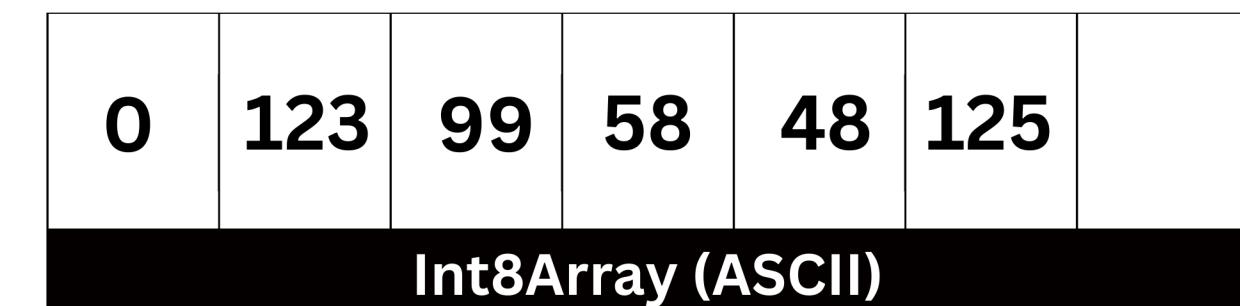
```
map.set(`c`, 0);
```

SharedArrayBuffer

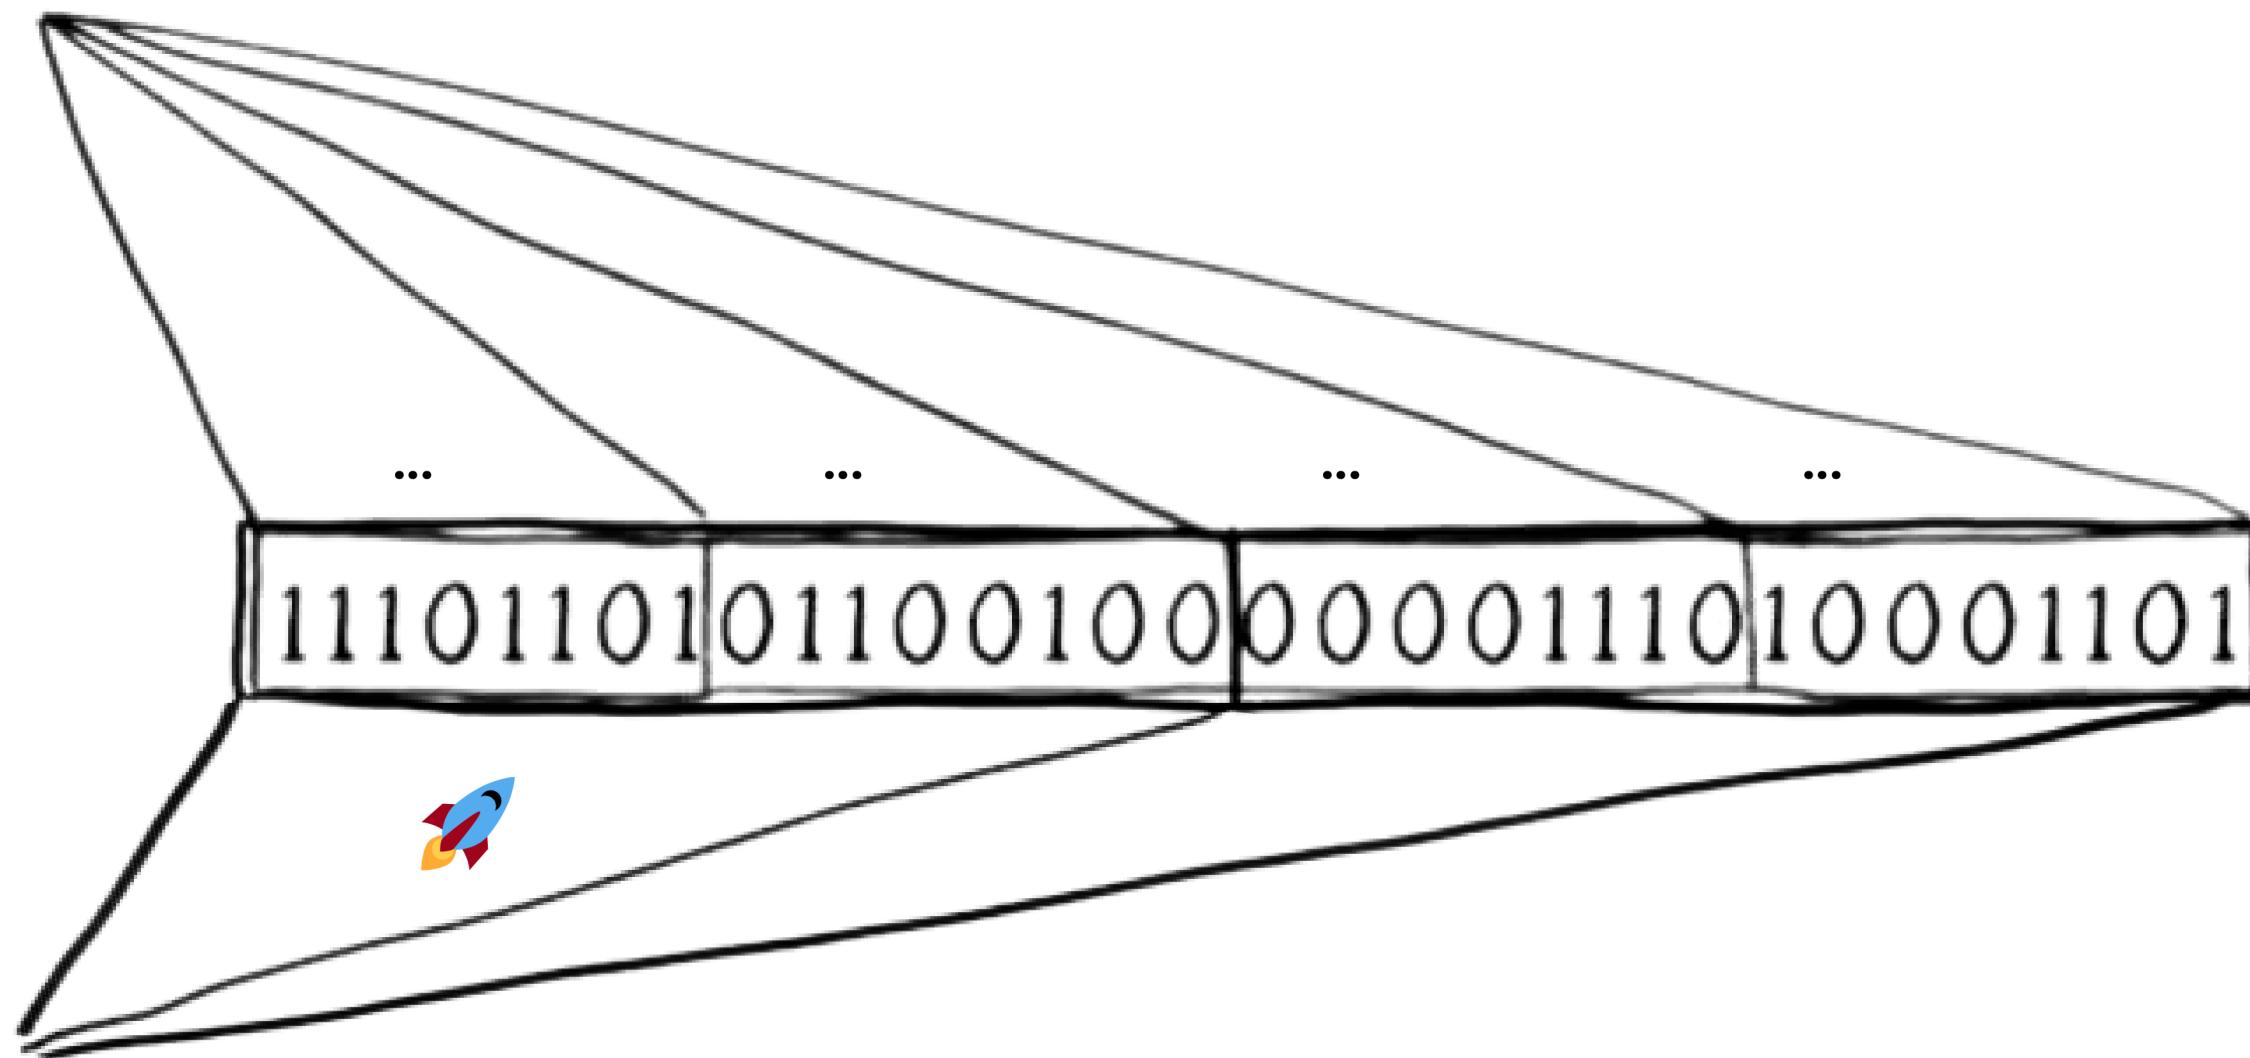


```
map.get(`c`);
```

0 - unlocked
1 - locked



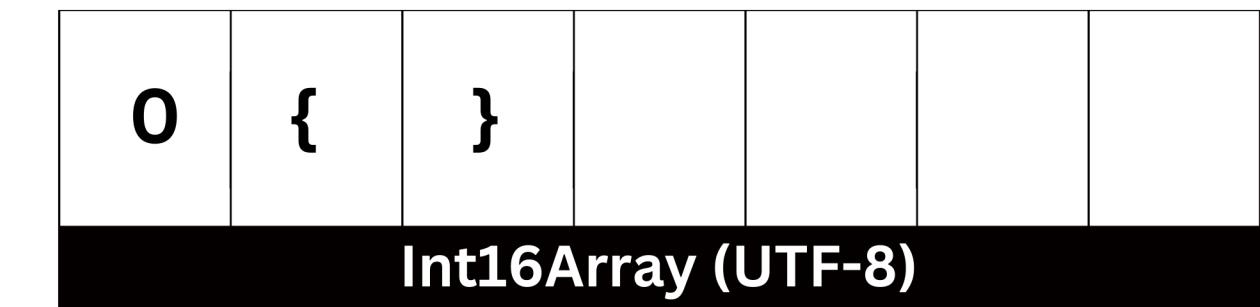
Int8Array



Int16Array

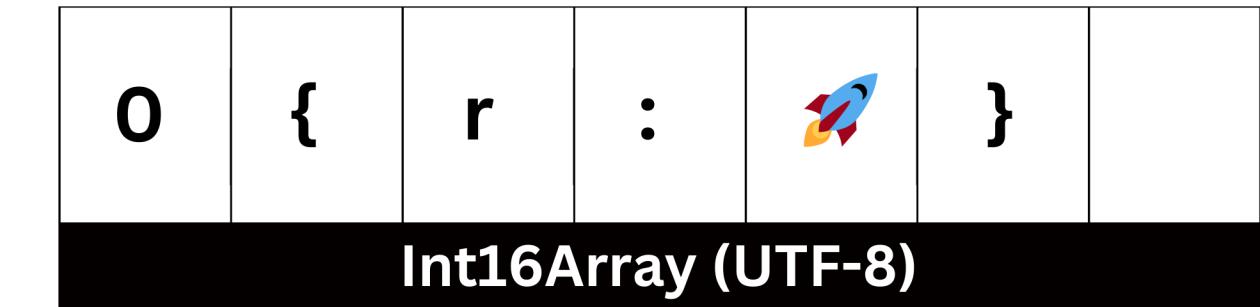
```
const map = new WorkerMap();
```

SharedArrayBuffer (64 bytes)



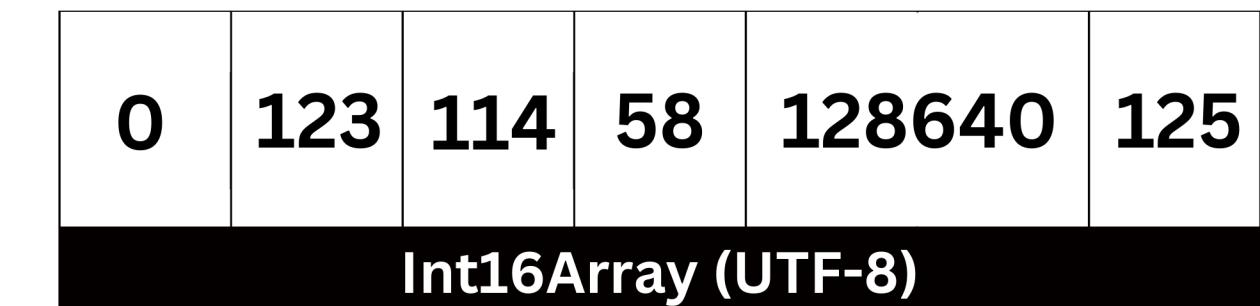
```
map.set('r', 🚀);
```

SharedArrayBuffer (64 bytes)



```
map.get('r');
```

0 - unlocked
1 - locked



SharedArrayBuffer

▼ Constructor

SharedArrayBuffer() constructor

▼ Properties

SharedArrayBuffer[@@species]

SharedArrayBuffer.prototype.byteLeng

SharedArrayBuffer.prototype.growable

SharedArrayBuffer.prototype.maxByteL

▼ Methods

SharedArrayBuffer.prototype.grow()

SharedArrayBuffer.prototype.slice()

Growing SharedArrayBuffers

SharedArrayBuffer objects can be made **growable** by including the `maxByteLength` option when calling the [`SharedArrayBuffer\(\)`](#) constructor. You can query whether a SharedArrayBuffer is **growable** and what its maximum size is by accessing its `growable` and `maxByteLength` properties, respectively. You can assign a new size to a **growable** SharedArrayBuffer with a [`grow\(\)`](#) call. New bytes are initialized to 0.

These features make growing SharedArrayBuffer s more efficient — otherwise, you have to make a copy of the buffer with a new size. It also gives JavaScript parity with WebAssembly in this regard (Wasm linear memory can be resized with [`WebAssembly.Memory.prototype.grow\(\)`](#)).

For security reasons, SharedArrayBuffer s cannot be reduced in size, only grown.

Available in Node >= v20

Is the `worker_thread`
lightweight?

Any questions





THANK YOU!



github.com/nairihar/ijs-conf-2023-munich

@nairihar



@nairihar